
On Variance Reduction in Learning Mean Flows

Anonymous Author(s)

Affiliation

Address

email

Abstract

One-step generative modeling has emerged as a leading approach to amortize the inference cost of diffusion and flow-matching models. Among distillation-free methods, MeanFlow training is notoriously unstable, with non-decreasing loss and unbounded gradient variance. In this work, we establish a theory that attributes this pathology to a misuse of the conditional velocity field: it plays two distinct statistical roles in the loss, both as an unbiased regression target and as a Monte Carlo control variate inside a Jacobi-vector product, with the original loss assigning the wrong coefficient to the latter. We derive the optimal coefficient in closed form, and show that a family of fixes in concurrent works corresponds to different practical realizations of the same optimum. A controlled sweep of this coefficient on two-dimensional benchmarks and on a latent Diffusion Transformer recovers the predicted bias-variance ordering. The optimal coefficient yields up to a 54% improvement in sample quality on two-dimensional benchmarks and a monotone FID trend at every matched-step DiT checkpoint. Crucially, the same DiT measurement also reveals a quantitative *FID-MSE landscape mismatch*: although gradient variance is minimized at an interior coefficient value, the coefficient that minimizes FID prefers the direct use of conditional velocity.

1 Introduction

Deep generative models that leverage neural networks to parametrize and sample from unknown data distributions have achieved remarkable success [1, 2, 3, 4, 5]. Among these models, flow-based generative models [6, 7, 8, 9, 10, 11, 12] learn a continuous-time transformation between a simple prior and the data distribution. Conditional flow matching [10, 11] and stochastic interpolants [12, 13] further enable simulation-free training by regressing a velocity field, achieving sample quality competitive with diffusion models [1, 14, 2]. However, integration at inference time remains a challenge due to its computational costs and numerical instability, motivating works on *one-step generation*, including progressive distillation [15], distribution matching [16, 17], consistency models [18, 19], and flow map matching [20]. Nevertheless, most of them require a pre-trained distillation teacher.

Mean Flows [21] provide a *distillation-free* alternative by learning a two-parameter average velocity field $\mathbf{u}(\mathbf{x}, r, t)$ through regressing a total-derivative identity between the average and instantaneous velocity. In principle, this identity yields exact self-supervision. In practice, however, training is plagued by *non-decreasing loss* and prohibitively *high gradient variance* [22, 23]. A plethora of concurrent works have proposed remedies that target different mechanisms. AlphaFlow [22] attributes slow convergence to gradient conflict between flow-matching and total-derivative consistency components and proposes an α -curriculum. Improved MeanFlow [23] replaces the conditional-velocity tangent with a learned marginal-velocity head. Re-MeanFlow [24] preprocesses the data with rectified-flow straightening to reduce trajectory curvature. Terminal Velocity Matching [25] sidesteps the spatial Jacobian by differentiating with respect to the terminal time. Functional Mean Flows [26] extend MeanFlow to Hilbert spaces and study marginal-

conditional consistency for functional data. Each fix is empirically effective in isolation, but there is no theory that explains *why* the original objective is unstable or what these fixes have in common. In this paper, we establish our theory around a single research question:

Why do stochastic tangents destabilize learning Mean Flows?

We answer this question through the lens of *variance reduction*. Specifically, we identify that the conditional velocity $\mathbf{v}_{\text{cond}} = \mathbf{x}_1 - \mathbf{x}_0$ plays two statistically distinct roles in the MeanFlow loss. As the regression target, it is an *unbiased* estimator of the marginal velocity $\mathbf{v}$. As the tangent inside the Jacobian-vector product (JVP) of the total-derivative identity, however, the same $\mathbf{v}_{\text{cond}}$ acts as a Monte Carlo *control variate* for the inaccessible marginal velocity, with a statistically suboptimal coefficient. The bias-variance trade-off generated by this substitution governs gradient variance through the quadratic form $\mathbf{g}^\top \mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top \mathbf{g}$, where $\mathbf{g} = \nabla_{\boldsymbol{\theta}} \mathbf{u}_{\boldsymbol{\theta}}$ and $\mathbf{J}$ is a spatial-Jacobi factor. The conditional fluctuation $\mathbf{v}_{\text{cond}} - \mathbf{v}$ is amplified by $\mathbf{J}$ at each gradient step, and stop-gradient prevents this amplification from reaching the optimizer, leading to the empirical pathology.

Following the theory, we derive the optimal tangent control-variate coefficient in closed form and show that it converges to a deterministic-tangent estimator in the practical regime where a marginal-velocity estimator is available. This result *explains* why deterministic tangents (*e.g.*, *EMA teachers and learned velocity heads*) all stabilize training: they are different practical instantiations of the same statistical optimum. To validate the framework empirically, we instantiate the $\beta = 1$ corner of β^* via an EMA-derived proxy with a flow-matching anchor (one extra JVP per step relative to vanilla MeanFlow), and run a controlled β -sweep ($\beta \in \{0, 0.25, 0.5, 1\}$) at toy and DiT-B/4 / ImageNet-256 scale. Direct measurement of per-step gradient variance recovers the predicted non-decreasing loss component $\text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top)$ and a $1.7 \times \sim 4.7 \times$ reduction at $\beta = 1$ on a 2-D variance-probing sub-sweep. At DiT scale the same sweep also exposes a quantitative *FID-vs-MSE landscape mismatch*: the FID ordering matches the bias-variance prediction across all four β values, but the FID-axis offset ratio across $\beta = 0.5$ vs $\beta = 1$ is super-linear in the MSE-axis bias coefficient, so the FID-optimal β shifts past the gradient-MSE interior minimizer to the unbiased corner. Our contributions in this paper are as follows:

- C1** We prove that per-step gradient variance in the original MeanFlow scales unboundedly with $\mathbf{g}^\top \mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top \mathbf{g}$ and the stop-gradient operator induces a *semi-gradient gap* that hides this term from the optimizer.
- C2** We frame the JVP tangent as a control variate for the marginal velocity, derive the optimal coefficient β^* in closed form, and unify concurrent works by showing that they each correspond to a different practical realization of β^* .
- C3** We empirically validate the framework via a controlled β -sweep at both 2-D and DiT-B/4 / ImageNet-256 scale, instantiating the $\beta = 1$ corner with an EMA proxy and flow-matching anchor. We observe $1.7 \times \sim 4.7 \times$ direct gradient-variance reduction, up to 54% SW_1 improvements on toy benchmarks, the predicted bias-variance FID ordering at every matched-step DiT checkpoint after the early-training transient, and a quantitative FID-vs-MSE landscape mismatch: the gradient-MSE optimum sits near $\beta \approx 0.94$ (directly measured matrix-form; scalar-isotropic prediction $\beta^* \in [0.46, 0.50]$) while the FID optimum lies at the unbiased corner $\beta = 0$, with the empirical $\Delta \text{FID}(\beta = 1)$ exceeding the MSE prediction by $\sim 7 \times$.

2 Preliminaries

Flow-Matching. Let $\mathbf{x}$ denote a data sample in $\mathcal{X} \subset \mathbb{R}^d$. Flow-matching generative models [10, 11] learn a time-dependent diffeomorphism $\psi(\mathbf{x}, t) : \mathbb{R}^d \times [0, 1] \rightarrow \mathbb{R}^d$ from a simple distribution $\mathbf{x}_1 \sim p(\mathbf{x}, 1)$ to the target $\mathbf{x}_0 \sim p(\mathbf{x}, 0) = p_{\text{data}}$ by parameterizing a velocity field $\mathbf{v}(\mathbf{x}, t) \triangleq \frac{d}{dt} \psi(\mathbf{x}, t)$. Since the exact *marginal velocity field* $\mathbf{v}(\mathbf{x}, t)$ is generally inaccessible, conditional flow matching instead learns a conditional field $\mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0) = \mathbf{x}_1 - \mathbf{x}_0$ with $\mathbf{x}_t = (1 - t)\mathbf{x}_0 + t\mathbf{x}_1$. The following lemma justifies the substitution. See Appendix C for its proof.

Lemma 1. *Given vector fields $\mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0)$ generating probability paths $p(\mathbf{x}, t \mid \mathbf{x}_0)$, for any $\mathbf{x}_0 \sim p(\mathbf{x}_0)$, the expected conditional velocity field is equal to the marginal velocity field*

$$\mathbf{v}(\mathbf{x}, t) = \mathbb{E}_{\mathbf{x}_0 \sim p(\mathbf{x}_0) \mid \mathbf{x}_t = \mathbf{x}}[\mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0)]. \quad (1)$$

89 **One-step Mean Flows.** Iterative integration of $\mathbf{v}(\mathbf{x}, t)$ at inference is expensive and can be numer-
 90 ically unstable. To bypass it, Mean Flows [21] learn a two-parameter average velocity field $\mathbf{u}(\mathbf{x}, r, t)$
 91 through regressing the identity along the flow trajectory $\mathbf{x}_\tau = \psi(\mathbf{x}_r, \tau)$ initialized at $\mathbf{x}_r = \mathbf{x}$:

$$(t - r) \mathbf{u}(\mathbf{x}, r, t) = \int_r^t \mathbf{v}(\mathbf{x}_\tau, \tau) d\tau. \quad (2)$$

92 Differentiating with respect to the terminal time t along the same trajectory yields

$$\mathbf{u}(\mathbf{x}, r, t) + (t - r) \frac{d}{dt} \mathbf{u}(\mathbf{x}, r, t) = \mathbf{v}(\mathbf{x}, t), \quad \frac{d}{dt} \mathbf{u} = \partial_{\mathbf{x}} \mathbf{u} \cdot \mathbf{v}(\mathbf{x}, t) + \partial_t \mathbf{u}, \quad (3)$$

93 where the total derivative $\frac{d}{dt} \mathbf{u}$ is essentially a JVP, denoted by $\text{JVP}(\mathbf{u}, (\mathbf{x}, r, t), (\mathbf{v}(\mathbf{x}, t), 0, 1))$. The
 94 original MeanFlow paper [21] replaces *both* occurrences of the marginal field with the conditional
 95 field $\mathbf{v}_{\text{cond}} = \mathbf{x}_1 - \mathbf{x}_0$, giving the MeanFlow loss

$$\mathcal{L}_{\text{MF}}(\boldsymbol{\theta}) = \mathbb{E}_{r, t, \mathbf{x}_0, \mathbf{x}_1} \left[\left\| \mathbf{u}_{\boldsymbol{\theta}} + (t - r) \text{sg}[\text{JVP}(\mathbf{u}_{\boldsymbol{\theta}}, (\mathbf{z}, r, t), (\mathbf{v}_{\text{cond}}, 0, 1))] - \mathbf{v}_{\text{cond}} \right\|_2^2 \right], \quad (4)$$

96 with $r, t \sim \mathcal{U}(0, 1)$, $r \leq t$, $\mathbf{x}_0 \sim p_{\text{data}}$, $\mathbf{x}_1 \sim \mathcal{N}(0, \mathbf{I}_d)$, $\mathbf{z} = (1 - t)\mathbf{x}_0 + t\mathbf{x}_1$, and the stop-gradient
 97 $\text{sg}[\cdot]$ avoiding double backpropagation through the JVP.

98 3 Variance Reduction in Learning Mean Flows

99 The MeanFlow loss in equation 4 is empirically known to be *non-decreasing* and to suffer from
 100 *high-variance gradients* [22, 23]. We investigate the statistical mechanism behind this pathology.
 101 Importantly, our analysis isolates the role of the conditional velocity field $\mathbf{v}_{\text{cond}}$ when it appears
 102 as the JVP tangent, and identifies that it acts as a Monte Carlo control variate for the inaccessible
 103 marginal velocity, with a coefficient that the original MeanFlow fixes to a suboptimal value. We
 104 derive the optimal coefficient and establish a connection to practical fixes in concurrent works.

105 3.1 Two roles of the conditional velocity field

106 With Reynolds decomposition, we express the conditional velocity by $\mathbf{v}_{\text{cond}} = \mathbf{v}(\mathbf{x}_t, t) + \mathbf{v}'$ with
 107 $\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t}[\mathbf{v}'] = \mathbf{0}$ according to equation 1 and $\Sigma_{\mathbf{v}'} \triangleq \text{Cov}_{\mathbf{x}_0 | \mathbf{x}_t}[\mathbf{v}']$. The MeanFlow loss in equation 4
 108 substitutes $\mathbf{v}_{\text{cond}}$ in two distinct positions: as the *regression target* $-\mathbf{v}_{\text{cond}}$, and as the *tangent* inside
 109 the total derivative. Carrying the decomposition through per-sample loss $\ell_{\text{MF}}(\boldsymbol{\theta})$ yields:

$$\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t}[\ell_{\text{MF}}(\boldsymbol{\theta})] = \|\mathbf{r}_{\boldsymbol{\theta}}^{\text{sg}}\|_2^2 + \text{sg}[\text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top)], \quad (5)$$

110 where $\mathbf{J} \triangleq (t - r) \partial_{\mathbf{z}} \mathbf{u}_{\boldsymbol{\theta}} - \mathbf{I}_d \in \mathbb{R}^{d \times d}$ is a *Jacobi factor* determined by the Jacobian matrix $\partial_{\mathbf{z}} \mathbf{u}_{\boldsymbol{\theta}}$
 111 and $\mathbf{r}_{\boldsymbol{\theta}}^{\text{sg}} \triangleq \mathbf{u}_{\boldsymbol{\theta}} + (t - r) \text{sg}[\partial_{\mathbf{z}} \mathbf{u}_{\boldsymbol{\theta}} \cdot \mathbf{v} + \partial_t \mathbf{u}_{\boldsymbol{\theta}}] - \mathbf{v}$ is the deterministic mean-field residual. Since the
 112 residual $\mathbf{r}_{\boldsymbol{\theta}}^{\text{sg}}$ is the exact loss we aim to minimize given by equation 3, the trace term leads to the
 113 non-decreasing loss and eliminates *iff* the Jacobi satisfies $\partial_{\mathbf{z}} \mathbf{u}_{\boldsymbol{\theta}} = \frac{1}{t - r} \mathbf{I}_d$ (equivalently, $\mathbf{J} = \mathbf{0}$). As
 114 a result, we identify two roles of the conditional velocity field in this decomposition:

115 **Unbiased Target.** As the regression target, $\mathbf{v}_{\text{cond}}$ injects $\mathbf{v}'$ *linearly* into the residual, contributing
 116 the *irreducible* per-sample noise floor $\text{Tr}(\Sigma_{\mathbf{v}'})$. This contribution is independent of $\boldsymbol{\theta}$ and reflects
 117 the natural Monte Carlo cost of replacing the inaccessible $\mathbf{v}$ by a per-sample proxy.

118 **Variance Amplifier.** As the JVP tangent, the same $\mathbf{v}_{\text{cond}}$ injects $\mathbf{v}'$ *multiplicatively* through the Ja-
 119 cobi factor $\mathbf{J}$, contributing the *amplified* term $\text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top)$. It scales unboundedly with the spectral
 120 magnitude of $\mathbf{J}$. We identify this amplification as the dominant driver of variance at the gradient
 121 level:

122 **Theorem 2** (Jacobian Variance Amplification). *Let $\mathbf{g} \triangleq \nabla_{\boldsymbol{\theta}} \mathbf{u}_{\boldsymbol{\theta}}(\mathbf{x}_t, r, t) \in \mathbb{R}^{d \times p}$ be the param-*
 123 *eter Jacobian of the average velocity. The trace of the conditional gradient covariance of $\ell_{\text{MF}}(\boldsymbol{\theta})$*
 124 *in equation 4 satisfies*

$$\text{Tr}(\text{Cov}[\nabla_{\boldsymbol{\theta}} \ell_{\text{MF}} | \mathbf{x}_t]) \propto \text{Tr}(\mathbf{g}^\top \mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top \mathbf{g}).$$




Figure 1: Spatial heatmap of the RMS conditional fluctuation $\sqrt{\text{Tr}(\Sigma_{v'} | \mathbf{x}_t)} = \sqrt{\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t} \|\mathbf{v}'\|^2}$ at three timesteps on a 2-D three-mode mixture. The norm is taken *before* the expectation, so unlike $\mathbb{E}[\mathbf{v}'] = \mathbf{0}$ this quantity is non-zero; it is the second moment that appears in Theorem 2 and concentrates in mode-mixing regions.

See the proof in Appendix E. Our Theorem 2 identifies that the scalar gradient variance grows with $\|\mathbf{J}\|^2$ (for fixed $\mathbf{g}$ and $\Sigma_{v'}$), saturates only at the irreducible noise floor when $\mathbf{J} \rightarrow \mathbf{0}$, and is *uncontrolled* by any term in the loss due to the use of the stop-gradient in the original MeanFlow. We notice that, without the stop-gradient operator, the optimizer could in principle reduce variance by driving $\mathbf{J} \rightarrow \mathbf{0}$. This establishes a *semi-gradient gap*, formally,

Theorem 3 (Semi-Gradient Gap). *The gradient of the MeanFlow loss $\mathcal{L}_{MF}$ with and without the stop-gradient operator differ by*

$$\underbrace{2(t-r) \mathbb{E} \left[\left(\nabla_{\theta} (\partial_z \mathbf{u}_{\theta} \cdot \mathbf{v} + \partial_t \mathbf{u}_{\theta}) \right)^{\top} \mathbf{r}_{\theta} \right]}_{\text{mean-field gradient difference}} + \underbrace{\mathbb{E} \left[\nabla_{\theta} \text{Tr}(\mathbf{J} \Sigma_{v'} \mathbf{J}^{\top}) \right]}_{\text{variance-driven gradient difference}}.$$

At convergence the mean-field difference vanishes ($\mathbf{r}_{\theta} \rightarrow \mathbf{0}$) and the variance-driven term $\text{Tr}(\mathbf{J} \Sigma_{v'} \mathbf{J}^{\top})$ dominates, but the stop-gradient hides it from the optimizer, producing the empirical non-decreasing-loss pathology. Figure 1 visualizes the per-state RMS fluctuation $\sqrt{\text{Tr}(\Sigma_{v'} | \mathbf{x}_t)}$ on a 2-D Gaussian mixture: the magnitude is non-zero, concentrates in mode-mixing regions where conditional paths overlap, and grows toward the latent endpoint.

3.2 Rethinking tangent as a control variate

The amplification in Theorem 2 arises due to using $\mathbf{v}_{\text{cond}} = \mathbf{v} + \mathbf{v}'$ as the JVP tangent injects the zero-mean fluctuation $\mathbf{v}'$. This actually corresponds to the canonical structure of a Monte Carlo *control variate* [27], where $\mathbf{v}_{\text{cond}}$ is a stochastic estimator of the inaccessible $\mathbf{v}$ with known mean, and the coefficient on the noise term $\mathbf{v}'$ can be a free parameter that determines the bias-variance trade-off. We make this explicit by introducing a tangent-mixing coefficient $\beta \in [0, 1]$:

$$\tilde{\mathbf{v}}_{\text{tang}}(\beta) \triangleq (1 - \beta) \mathbf{v}_{\text{cond}} + \beta \hat{\mathbf{v}} = \mathbf{v} + (1 - \beta) \mathbf{v}' + \beta \mathbf{b}, \quad (6)$$

where $\hat{\mathbf{v}}$ can be any deterministic proxy for the marginal velocity (e.g., a *velocity network*) with bias $\mathbf{b} \triangleq \hat{\mathbf{v}} - \mathbf{v}$ that is deterministic given $\mathbf{x}_t$. Herein, $\beta = 0$ recovers vanilla MeanFlow and $\beta = 1$ replaces the tangent entirely with the deterministic proxy.

Bias-variance trade-off. Substituting $\tilde{\mathbf{v}}_{\text{tang}}(\beta)$ for $\mathbf{v}_{\text{cond}}$ in the JVP tangent and propagating through the loss, the per-sample gradient writes

$$\mathbf{g}^{(\beta)} = 2\mathbf{g}^{\top} \left[\mathbf{r}_{\theta} + ((1 - \beta) \mathbf{J} - \beta \mathbf{I}) \mathbf{v}' + \beta (\mathbf{J} + \mathbf{I}) \mathbf{b} \right], \quad (7)$$

where the noise term scales the conditional fluctuation $\mathbf{v}'$ by the matrix $((1 - \beta) \mathbf{J} - \beta \mathbf{I})$, and the bias term scales the proxy error $\mathbf{b}$ by $\beta (\mathbf{J} + \mathbf{I})$. This establishes the following theorem.

Theorem 4 (Optimal control-variate coefficient). *Define the conditional MSE of the per-sample gradient equation 7 as its mean-squared error against its noise- and bias-free target $2\mathbf{g}^{\top} \mathbf{r}_{\theta}$, taken over the randomness of $\mathbf{v}'$ at fixed $\mathbf{x}_t$:*

$$M(\beta) \triangleq \mathbb{E}_{\mathbf{v}' | \mathbf{x}_t} [\|\mathbf{g}^{(\beta)} - 2\mathbf{g}^{\top} \mathbf{r}_{\theta}\|_2^2].$$

Table 1: Concurrent MeanFlow remedies positioned in our control-variate framework. Each method either drives $\beta \rightarrow 1$ with a deterministic proxy, reduces the amplification factor $\|\mathbf{J}\|^2$ directly, or replaces the JVP construction altogether.

Method	Mechanism in our framework	Practical realization
AlphaFlow [22]	avoid high- κ regimes	α -curriculum schedule
Improved MF [23]	deterministic tangent ($\beta \rightarrow 1$)	separate velocity head + sg
Modular MF [28]	deterministic tangent ($\beta \rightarrow 1$)	gradient-modulated loss family
Kim <i>et al.</i> [29]	deterministic tangent ($\beta \rightarrow 1$)	staged training schedule
TVM [25]	remove the spatial Jacobian ($\mathbf{J} \rightarrow \mathbf{0}$)	terminal-time differentiation
Re-MeanFlow [24]	reduce $\ \mathbf{J}\ $	rectified-flow preprocessing
Decoupled MF [30]	bypass JVP construction	flow-map conditioning on later t
$\beta = 1$ (ours)	deterministic tangent ($\beta \rightarrow 1$)	EMA proxy + FM anchor

Under the scalar-isotropic approximation $\Sigma_{\mathbf{v}'} \approx \sigma^2 \mathbf{I}_d$, $\mathbf{J} \approx \kappa \mathbf{I}_d$, and the parameter-isotropy approximation $\mathbf{g}\mathbf{g}^\top \propto \mathbf{I}_d$ (which reduces M to a per-component MSE; cf. Appendix E),

$$M(\beta) \propto \beta^2(\kappa+1)^2\|\mathbf{b}\|^2 + \sigma^2 d((1-\beta)\kappa - \beta)^2.$$

For $\kappa > 0$, $M(\beta)$ admits a unique minimizer

$$\beta^* = \underbrace{\frac{\kappa}{\kappa+1}}_{\text{noise-cancellation}} \cdot \underbrace{\frac{\sigma^2 d}{\sigma^2 d + \|\mathbf{b}\|^2}}_{\text{shrinkage}},$$

with optimum value $M(\beta^*) \propto \sigma^2 d \kappa^2 \|\mathbf{b}\|^2 / (\sigma^2 d + \|\mathbf{b}\|^2)$. For all κ , $\|\mathbf{b}\|^2 > 0$, $M(\beta^*)$ strictly dominates both corner MSEs $M(0) \propto \sigma^2 \kappa^2 d$ (vanilla MeanFlow) and $M(1) \propto (\kappa+1)^2 \|\mathbf{b}\|^2 + \sigma^2 d$ (deterministic-tangent estimator with bias). The boundary case $\kappa = -1$ (which corresponds to $\partial_{\mathbf{z}} \mathbf{u}_\theta = \mathbf{0}$ at initialization) leaves $M(\beta)$ constant and is excluded.

The proof is a direct application of bias-variance decomposition to equation 7 and is given in Appendix E. The two factors of β^* have crisp interpretations: $\kappa/(\kappa+1)$ is the coefficient that *cancels the Jacobian-amplified noise* via destructive interference between the $\mathbf{J}\mathbf{v}'$ and $\mathbf{I}\mathbf{v}'$ contributions in equation 7; $\sigma^2 d / (\sigma^2 d + \|\mathbf{b}\|^2)$ is the standard James-Stein shrinkage that pulls β^* toward the unbiased corner when the proxy bias $\|\mathbf{b}\|$ is large. As the proxy bias shrinks ($\|\mathbf{b}\|^2 \rightarrow 0$), $\beta^* \rightarrow \kappa/(\kappa+1)$, and as the Jacobian magnitude grows ($\kappa \rightarrow \infty$), $\beta^* \rightarrow 1$.

Vanilla MeanFlow is generally suboptimal. Vanilla MeanFlow uses $\beta = 0$ which is optimal only when $\kappa = 0$ (i.e., $\partial_{\mathbf{z}} \mathbf{u}_\theta = \frac{1}{t-r} \mathbf{I}_d$) or when no deterministic proxy is available ($\|\mathbf{b}\| = \infty$). Both conditions fail in practice: trained backbones produce $\partial_{\mathbf{z}} \mathbf{u}_\theta$ matrices that are spectrally far from $\frac{1}{t-r} \mathbf{I}_d$, and an EMA copy of $\mathbf{u}_\theta$ provides a viable deterministic proxy at near-zero compute cost.

Concurrent fixes through this lens. Several remedies for MeanFlow’s training pathology in concurrent work can now be located in the same one-parameter family. Each fix corresponds to a different practical realization of the optimum, distinguished by *which proxy* they use for $\mathbf{v}$ and *which value of β* the construction implicitly selects:

- **Deterministic-tangent methods** (Improved MF [23], EMA tangents) take $\beta \rightarrow 1$ with a learned or EMA-derived proxy $\hat{\mathbf{v}}$. The amplified noise term is eliminated; the bias term $(\mathbf{J} + \mathbf{I})\mathbf{b}$ is controlled by training the proxy.
- **Reparameterization methods** (TVM [25]) sidestep the spatial Jacobian by reparameterizing the total derivative, effectively setting $\mathbf{J} = \mathbf{0}$ in our framework.
- **Curvature-reduction methods** (Re-MeanFlow [24]) reduce $\|\mathbf{J}\|$ via reflow preprocessing, lowering κ and the amplified-noise term.
- **Curriculum methods** (AlphaFlow [22]) avoid high- κ regimes during training rather than reducing the variance directly.

None of these methods explicitly identify the control-variate structure or derive β^* . We do, and we use the result both as an analytical lens for these methods and as the foundation for our practical training framework.

3.3 Empirical instantiation of the $\beta = 1$ corner

To validate the framework empirically, we need a concrete realization of β^* that we can sweep at scale. We construct a controlled instance of the $\beta = 1$ corner of Theorem 4: the JVP tangent is replaced by an EMA-derived deterministic proxy $\hat{v} = \mathbf{u}_{\hat{\theta}}(\mathbf{x}_t, t, t)$, and a small flow-matching anchor at $r = t$ keeps the proxy bias $\mathbf{b} = \hat{v} - \mathbf{v}$ small so the bias term $\beta(\mathbf{J} + \mathbf{I})\mathbf{b}$ in equation 7 stays controlled. The regression target remains $\mathbf{v}_{\text{cond}}$ unchanged; only the tangent is replaced, exploiting the role asymmetry of Section 3.1. The construction adds one extra JVP per step relative to vanilla MeanFlow. Appendix D gives the full loss, training algorithm, and three propositions characterizing the resulting gradient form, residual bias, and the necessity of the target/tangent role split. In Section 5 we sweep $\beta \in \{0, 0.25, 0.5, 1\}$ with the EMA-tangent recipe realizing the $\beta = 1$ corner, the linear interpolation realizing $\beta \in \{0.25, 0.5\}$, and vanilla MeanFlow realizing $\beta = 0$.

4 Related Work

MeanFlow and its training pathology. Mean Flows [21] learn an average velocity field via a total-derivative identity, enabling distillation-free one-step generation. The empirical pathology, non-decreasing loss with high gradient variance, was diagnosed in concurrent work from several angles. AlphaFlow [22] attributes slow convergence to gradient conflict between flow-matching and total-derivative consistency components and proposes an α -curriculum. Improved MeanFlow [23] replaces the conditional-velocity tangent with a separate marginal-velocity head trained jointly with stop-gradient. Kim *et al.* [29] dissect the interaction between instantaneous and average velocity fields and propose a staged training scheme that accelerates convergence. Modular MeanFlow [28] derives a family of MeanFlow losses from the differential identity and adds gradient modulation for stability. Re-MeanFlow [24] addresses trajectory curvature through reflow pre-processing that straightens conditional paths. Decoupled MeanFlow [30] converts a pretrained flow model into a flow map by conditioning the final blocks on a subsequent timestep, sidestepping joint training of the average velocity. Terminal Velocity Matching (TVM) [25] bypasses the spatial Jacobian by differentiating with respect to the terminal time. Functional Mean Flows [26] extends MeanFlow to Hilbert spaces; Riemannian MeanFlow [31] extends it to manifolds; and OT-MeanFlow [32] couples the construction with optimal transport for 3D point-cloud generation. Each of these works proposes an effective practical fix or domain extension; in contrast, our work isolates the statistical role of the JVP tangent itself, identifies the conditional fluctuation $\mathbf{v}'$ as a control variate with the wrong coefficient inside a semi-gradient JVP, and derives the optimal coefficient. Concurrent fixes that act on the tangent or proxy correspond to different practical realizations of this optimum (Table 1); curvature-reduction methods (Re-MeanFlow) instead shrink $\|\mathbf{J}\|$ in our framework, which lowers the amplification factor by Theorem 2.

Variance reduction in flow training. Several recent works address variance in flow-based training from a complementary angle. Stable Velocity [33] reveals a two-regime variance structure in flow matching and proposes composite conditioning; TPC [34] reduces variance through temporal pair coupling; and preconditioned flow matching [35] addresses ill-conditioning via anisotropic preconditioning. Bertrand *et al.* [36] argue that stochastic-target variance is negligible for standard flow matching, consistent with our analysis, since the problematic amplification we identify is specific to the MeanFlow JVP, not the flow-matching loss itself. Rectified flows [37, 38] reduce trajectory curvature by iterative straightening, which by Theorem 2 reduces $\|\mathbf{J}\|^2$ and hence the amplified noise term.

Control variates and target networks. The control-variate framework we apply has classical roots in Monte Carlo estimation [39]; it is also the analytical structure underlying target networks in deep RL [40] and self-distillation in self-supervised learning [41]. We make the connection explicit by interpreting the EMA-derived tangent in MeanFlow as a target-network-style proxy that drives the control-variate coefficient toward its optimum.

5 Empirical Validation of the Control-Variate Framework

We empirically validate the framework of Section 3 through a direct probe of the bias-variance landscape $M(\beta)$: rather than picking a single instantiation and benchmarking it against vanilla




Figure 2: Empirical gradient covariance trace $\text{Tr}(\text{Cov}[\nabla_{\theta} \ell_{\text{MF}}])(\beta)$ averaged over the second half of training, on four 2-D toy datasets ($\beta \in \{0, 0.25, 0.5, 0.75, 1\}$, single seed; gradient-variance probing has non-trivial overhead). Per-panel red annotations give the $\beta = 0$ over $\beta = 1$ reduction ratio. The monotone decrease in β on every dataset confirms Theorem 2’s prediction that the conditional-tangent fluctuation $\mathbf{v}'$ contributes the dominant gradient-variance term.

MeanFlow, we sweep the tangent-mixing coefficient $\beta \in [0, 1]$ end-to-end and compare the empirical landscape to the closed-form prediction of Theorem 4. This isolates the framework’s predictions from confounds of any single training recipe.

Setup. The interior- β instantiation equation 6 interpolates between vanilla MeanFlow ($\beta = 0$, conditional-velocity tangent) and the EMA-tangent corner ($\beta = 1$). For 2-D toy benchmarks we use a 3-layer MLP with 128 hidden units, 200,000 optimization steps, three random seeds $\{42, 0, 1\}$, and one shared backbone, sampling schedule, and optimizer per dataset; only β varies. We measure (i) the per-step gradient covariance $\text{Tr}(\text{Cov}[\nabla_{\theta} \ell_{\text{MF}}])$ from $K = 8$ replica mini-batches (Theorem 2) on a 4-dataset, single-seed sub-sweep at $\beta \in \{0, 0.25, 0.5, 0.75, 1\}$ (gradient-variance probing has non-trivial overhead), and (ii) the final-step sliced Wasserstein-1 (sample quality) on the full $\{0.0, 0.1, \dots, 1.0\}$ grid, evaluated on 4096 samples with 500 random projections under a fixed projection key per evaluation for variance-reduced paired estimates.

5.1 Variance reduction matches Theorem 2

Figure 2 reports $\text{Tr}(\text{Cov}[\nabla_{\theta} \ell_{\text{MF}}])(\beta)$ across the gradient-variance sub-sweep on four 2-D toy datasets, averaged over the second half of training to suppress transient. The relation is monotone decreasing in β on every dataset, with a $1.7\times \sim 4.7\times$ reduction at $\beta = 1$ relative to vanilla MeanFlow (per-dataset ratios in Figure 2). The largest reduction ($4.71\times$) appears on `eight_gaussians`, where the trained backbone has the largest spectral magnitude $\|\mathbf{J}\|$; the smallest ($1.73\times$) on `swiss_roll`, where the backbone is closer to the no-amplification regime $\partial_{\mathbf{z}} \mathbf{u}_{\theta} \approx \frac{1}{t-r} \mathbf{I}_d$. The empirical ordering tracks the variance term $\sigma^2 d((1-\beta)\kappa - \beta)^2$ from Theorem 4 and Theorem 2, providing direct empirical evidence that the conditional-tangent fluctuation $\mathbf{v}'$ is the dominant variance source at the gradient level.

5.2 The empirical $M(\beta)$ landscape

Whereas $\text{Tr}(\text{Cov}[\nabla_{\theta} \ell_{\text{MF}}])(\beta)$ probes the variance term in isolation, the sample-quality landscape $\text{SW}_1(\beta)$ exposes the full bias-variance trade-off. Figure 3 traces $\text{SW}_1(\beta)$ across the full 11-point grid on six 2-D toy datasets (3 seeds, error bars are SEM). On the five high-curvature datasets (`checkerboard`, `eight_gaussians`, `two_moons`, `swiss_roll`, `pinwheel`), the empirical β^* falls in $\{0.8, 1.0\}$, recovering the predicted near-corner location for the small- $\|\mathbf{b}\|^2$ regime where the EMA proxy is accurate (the James-Stein shrinkage factor $\sigma^2 d/(\sigma^2 d + \|\mathbf{b}\|^2)$ approaches 1). On the lower-curvature `two_spirals` dataset, the trained backbone produces $\partial_{\mathbf{z}} \mathbf{u}_{\theta}$ close to $\frac{1}{t-r} \mathbf{I}_d$, driving κ toward zero. The predicted β^* then shrinks toward the unbiased corner $\beta = 0$ through the noise-cancellation factor $\kappa/(\kappa+1)$ in Theorem 4, and the empirical sweep concurs ($\beta^* = 0$). On `swiss_roll` the SW_1 reduction at the empirical $\beta^* = 1$ relative to $\beta = 0$ is 51% ($0.098 \rightarrow 0.048$), and a parallel matched-method comparison of vanilla MeanFlow against the full $\beta = 1$ corner recipe (EMA tangent + FM anchor; Table 3) reaches 54%.




Figure 3: Empirical sample-quality landscape $SW_1(\beta)$ on six 2-D toy datasets ($\beta \in \{0.0, 0.1, \dots, 1.0\}$, three seeds, error bars are SEM). The red dotted line marks the empirical $\beta^* = \operatorname{argmin}_{\beta} SW_1(\beta)$ on each dataset. Five of six datasets settle at $\beta^* \in \{0.8, 1.0\}$ near the deterministic-tangent corner; `two_spirals` alone settles at $\beta^* = 0$ as the noise-cancellation factor $\kappa/(\kappa+1)$ shrinks toward 0 on this low-curvature dataset.

5.3 High-dimensional scaling on DGMM

We sweep $\beta \in \{0, 0.1, \dots, 1\}$ on a dense Gaussian mixture (DGMM) in $\mathbb{R}^d$, $d \in \{2, 4, 8, 16, 32, 64\}$ (three seeds $\{42, 0, 1\}$, 200k steps, otherwise identical setup to the 2-D toys; error bars are SEM). Figure 4 reports $SW_1(\beta)$ per dimension. Two regimes emerge: at low-to-moderate d the empirical β^* sits near the deterministic-tangent corner, while at higher d the optimum moves into the interior. The transition is consistent with Theorem 4: at low d the EMA proxy is highly accurate ($\|\mathbf{b}\|^2$ small, shrinkage $\sigma^2 d / (\sigma^2 d + \|\mathbf{b}\|^2) \approx 1$) and $\kappa/(\kappa+1)$ saturates near 1; at higher d the network’s Jacobian spectrum drifts off saturation, κ drops, and the noise-cancellation factor $\kappa/(\kappa+1)$ pulls β^* to the interior. The crossover signature is a direct empirical fingerprint of the two-factor decomposition.

5.4 Validation at DiT-B/4 / ImageNet-256 scale

To check that the framework’s predictions hold beyond the toy regime, we run a four-point sweep $\beta \in \{0, 0.25, 0.5, 1\}$ on a class-conditional DiT-B/4 [42] trained on Stable Diffusion VAE latents ($32 \times 32 \times 4$, $d = 4096$); see Appendix F for the full setup, FID table, diagnostic measurements, and DiT-specific bias and shrinkage trajectory.

Figure 5 visualizes the converged-FID landscape as a function of β , with empirical $\Delta FID(\beta) \triangleq FID(\beta) - FID_{\text{baseline}}$ (using each run’s final-logged FID) overlaid against the MSE-prediction curve $\beta^2 \cdot A$ (anchored at $\beta = 0.5$). The four-point matched-step ordering $FID(\beta = 0) < FID(\beta = 0.25) < FID(\beta = 0.5) < FID(\beta = 1)$ holds at every 5k-aligned checkpoint logged after the early-training transient (35 consecutive datapoints across all four runs from step 30k to step 200k; the $\beta = 0.25$ run is still in training, the rest are finished); per-checkpoint values are reported in Table 5.

FID-vs-MSE landscape mismatch. Direct measurement of both ingredients of Theorem 4 on the DiT checkpoints (EMA-tracking proxy $\|\mathbf{b}\|^2$, irreducible noise floor $\sigma^2 d$, both reported in Appendix F) places the formula-predicted optimum in the interior range $\beta^* \in [0.46, 0.50]$ under the scalar-isotropic approximation. The empirical $\beta = 0.5$ -beats- $\beta = 1$ ordering at convergence is consistent with that prediction. To test whether FID tracks the gradient-MSE landscape *quantitatively*, we anchor the MSE-prediction curve $\beta^2 \cdot A$ (with $A = \Delta FID(\beta = 0.5)/0.25$) and compare the $\beta = 1$ empirical offset to the curve. Empirically $\Delta FID(\beta = 1) \approx 11.24$ from the baseline floor




Figure 4: DGMM β -sweep in $\mathbb{R}^d$ for $d \in \{2, 4, 8, 16, 32, 64\}$ ($\beta \in \{0.0, 0.1, \dots, 1.0\}$, three seeds, error bars are SEM). Red dotted line marks the empirical $\beta^* = \arg \min_{\beta} SW_1(\beta)$ on each dimension. The optimum stays near the deterministic-tangent corner at low d and moves into the interior at the highest dimension, an empirical signature of $\kappa/(\kappa+1)$ coming off saturation as d grows.

301 FID_{baseline} = 12.12, while the MSE prediction at $\beta = 1$ is only 1.56 — a $\sim 7\times$ super-linear amplifica-
 302 tion on the FID axis (Figure 5). The FID landscape therefore over-penalizes bias relative to gradient
 303 MSE. As a consequence, the FID-optimal β shifts past the gradient-MSE interior minimizer all the
 304 way to the unbiased corner $\beta = 0$.

305 **Direct matrix-form measurement.** To verify that the scalar-isotropic prediction is not artifactual,
 306 we directly measured a matrix-form upper bound on β^*_{matrix} of Appendix E on the baseline checkpoint
 307 by sampling actual v' values from the conditional distribution (no Hutchinson probes; we average
 308 $(Jv') \cdot ((J+I)v')$ and $\|(J+I)v'\|^2$ across data samples, which is unbiased for the trace formulas
 309 regardless of how anisotropic $\Sigma_{v'}$ is). The estimator omits the bias term $\|(J+I)b\|^2$ from the
 310 denominator, giving

$$\beta^*_{\text{no bias}} \triangleq \text{Tr}(J\Sigma_{v'}(J+I)^\top) / \text{Tr}((J+I)\Sigma_{v'}(J+I)^\top) \geq |\beta^*_{\text{matrix}}|,$$

311 since the omitted bias term is non-negative. We probe at $t \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$, fixed gap
 312 $t - r = 0.25$, 512 samples per t (using the model’s own boundary $u_\theta(x_t, t, t)$ as the marginal-
 313 velocity proxy, asymptotically correct at $\beta = 0$ training convergence). The numerator is *strictly*
 314 *positive* at every probed t , ruling out the unbiased-corner regime $\beta^*_{\text{matrix}} \leq 0$. Aggregated across
 315 t , $\beta^*_{\text{no bias}} \approx 0.94$ — close to the deterministic-tangent corner and *larger* than the scalar-isotropic
 316 estimate $\beta^* \in [0.46, 0.50]$. The full β^*_{matrix} then lies in $(0, 0.94]$; we do not pin down the exact
 317 value because a clean estimator of the true model bias $b = u_\theta(x_t, t, t) - v(x_t, t)$ is intractable on
 318 ImageNet without access to the conditional distribution $p(x_0 | x_t)$, and substituting the EMA copy
 319 $u_{\bar{\theta}}$ for v would only measure an EMA-tracking gap rather than the true bias. The matrix-form
 320 measurement therefore strengthens, rather than weakens, the FID-vs-MSE mismatch: the gradient-
 321 MSE optimum sits in $(0, 0.94]$ while the FID optimum sits at the unbiased corner $\beta = 0$ — the strict
 322 interior-vs-corner gap on the gradient-MSE axis is preserved under the matrix-form analysis.

323 5.5 Discussion

324 The full-sweep validation establishes three points jointly. (i) *Variance scaling*: $\text{Tr}(\text{Cov}[\nabla_\theta \ell_{\text{MF}}])(\beta)$
 325 matches the variance term $\sigma^2 d((1-\beta)\kappa - \beta)^2$ of Theorem 4 on every dataset (Figure 2), with $1.7\times \sim$
 326 $4.7\times$ reduction at $\beta = 1$. Theorem 2 identifies the same Jacobi-amplified term as the variance source
 327 at the special case $\beta = 0$. (ii) *Bias-variance landscape*: the empirical $SW_1(\beta)$ landscape recovers the
 328 predicted shape on six 2-D toy datasets (Figure 3); the empirical β^* tracks the closed-form prediction

DiT-B/4 / ImageNet-256: FID-vs-MSE landscape mismatch




Figure 5: DiT-B/4 / ImageNet-256 FID-vs-MSE landscape mismatch. Empirical $\Delta\text{FID}(\beta) = \text{FID}(\beta) - \text{FID}_{\text{baseline}}$ at the converged (final-logged) step of each run; the dashed gray curve is the MSE-prediction $\beta^2 \cdot A$ anchored at $\beta = 0.5$. The $\beta = 1$ corner sits an order of magnitude above the MSE prediction (super-linear FID dependence on bias). The directly-measured matrix-form gradient-MSE optimum $\beta_{\text{matrix}}^* \approx 0.94$ (green tick on the x-axis) is far from the FID-optimal $\beta = 0$, visualizing the landscape mismatch. The $\beta = 0.25$ run is still in training and has a strict upper-bound marker.

329 within seed standard error, including the corner shrinkage to $\beta^* = 0$ on `two_spirals` where κ is
330 small. (iii) *Scale invariance*: at DiT-B/4 / ImageNet-256 ($d = 4096$) the predicted ordering across
331 $\beta \in \{0, 0.25, 0.5, 1\}$ holds at every matched-step checkpoint (Figure 5), and the empirical asymptotic
332 offsets reveal a quantitative super-linear FID dependence on the gradient-MSE-axis bias term —
333 corroborated by direct matrix-form measurement of $\beta^* \approx 0.94$ on the gradient-MSE axis versus
334 the FID-optimum at $\beta = 0$. Together these establish that the control-variate framework correctly
335 identifies both the gradient-variance source in MeanFlow training and the bias-variance ordering at
336 the loss-landscape level, across four orders of magnitude in dimensionality.

337 The framework also recasts a family of concurrent works (iMF [23], AlphaFlow [22], Re-
338 MeanFlow [24], TVM [25]) as practical realizations of the same one-parameter optimum at differing
339 β , κ , or $\|b\|^2$ regimes (Table 1). Our $\beta = 0.5$ DiT measurement shows that pushing $\beta \rightarrow 1$ pays a
340 quantifiable FID cost beyond what the gradient-MSE picture predicts. This observation, novel to
341 this work, suggests that practical β should be chosen jointly with the downstream evaluation metric
342 rather than uniformly at the gradient-MSE optimum.

343 6 Conclusion

344 We addressed the question of *why* stochastic tangents destabilize Mean Flow learning and provided
345 a unifying theoretical answer: the conditional velocity field plays two distinct statistical roles in
346 the MeanFlow loss (unbiased Monte Carlo target and Monte Carlo control variate inside a semi-
347 gradient JVP), and the original loss assigns the wrong coefficient to the latter. We derived the optimal
348 control-variate coefficient β^* in closed form and showed that vanilla MeanFlow’s choice ($\beta = 0$)
349 is generally suboptimal because per-step gradient variance scales unboundedly with $g^\top J \Sigma_v J^\top g$.
350 The stop-gradient operator induces a semi-gradient gap that hides this scaling from the optimizer,
351 producing the empirical non-decreasing-loss pathology. Our framework formally unifies a family
352 of concurrent fixes (Improved MF, TVM, Re-MeanFlow) as different practical realizations of the
353 optimum. To validate the framework empirically, we instantiated the $\beta = 1$ corner via an EMA
354 proxy with a flow-matching anchor and ran a controlled β -sweep ($\beta \in \{0, 0.25, 0.5, 1\}$) at toy and
355 DiT-B/4 / ImageNet-256 scale (Appendix F). The sweep recovers the predicted $1.7\times\sim 4.7\times$ gradient-
356 variance reduction with up to a 54% reduction in SW_1 on the high-curvature `swiss_roll` dataset,

and recovers the predicted bias-variance FID ordering at every matched-step DiT checkpoint after the early-training transient. The same DiT measurement also exposes a quantitative *FID-vs-MSE landscape mismatch*: the gradient-MSE optimum sits near $\beta \approx 0.94$ (matrix-form, directly probed; scalar-isotropic prediction $\beta^* \in [0.46, 0.50]$ is conservatively interior), while the FID optimum lies at the unbiased corner $\beta = 0$, and the empirical $\Delta\text{FID}(\beta = 1)$ exceeds the MSE-prediction (anchored at $\beta = 0.5$) by $\sim 7\times$. The FID landscape therefore over-penalizes bias relative to the gradient-MSE landscape on this benchmark, and a metric-aware β schedule is left for future work.

Practical guidance. The empirical β^* varies systematically with the dataset (via the data-dependent κ and $\|\mathbf{b}\|^2$) — exactly as Theorem 4 predicts. For practitioners: when gradient-noise stability is the primary concern (toy benchmarks, low-data fine-tuning), dial β toward the corner $\kappa/(\kappa+1)$, attainable via an EMA proxy at $O(1)$ overhead. When FID is the metric of record (large-scale image synthesis), our DiT measurement suggests dialing β back toward 0 or investing in interventions (longer EMA half-life, multi-EMA, distillation pretraining) that shrink $\|\mathbf{b}\|^2$ until the FID bias penalty becomes negligible. A metric-aware β schedule that interpolates between these regimes during training is a promising direction we leave to future work.

Limitations. (i) *Scalar-isotropic approximation.* Theorem 4 is derived under $\mathbf{J} \approx \kappa \mathbf{I}_d$ and $\Sigma_{\mathbf{v}'} \approx \sigma^2 \mathbf{I}_d$. In the full diagonal matrix-valued setting the optimum is direction-dependent: a per-direction $\beta_i = \frac{\mathbf{J}_{ii}}{\mathbf{J}_{ii}+1} \cdot \frac{\sigma_i^2}{\sigma_i^2+b_i^2}$ would minimize the per-component gradient MSE, and a single global β trades off residual variance across directions. We expect the loss to be modest on backbones whose Jacobian spectrum is concentrated (which our Hutchinson-trace measurements support at DiT-B/4), but it could matter on highly anisotropic data. (ii) *Single-seed DiT.* Our DiT-B/4 / ImageNet-256 evaluation is single-seed; multi-seed FID confirmation will be added for camera-ready. (iii) *FID-aware loss design.* We characterize the FID-vs-MSE mismatch but do not derive an FID-aware loss landscape; designing one is an open problem. (iv) *Non-Euclidean data.* The framework is stated for $\mathbb{R}^d$; extending the closed-form β^* to manifold MeanFlow [31] is straightforward at the gradient level but requires care in defining the second moment $\text{Tr}(\Sigma_{\mathbf{v}'})$ on tangent spaces.

Future work. Three directions stand out. (i) *Beyond MeanFlow.* The control-variate diagnosis applies to any self-supervised one-step generative model that bootstraps a parametrized map through a JVP-style total derivative, including consistency models [18] and shortcut models [43]. (ii) *Adaptive β .* A per-sample scalar coefficient driven by online Hutchinson estimates of κ and $\|\mathbf{b}\|^2$ would close the global-vs-per-sample scalar gap; closing the scalar-vs-matrix gap of (i) further requires direction-dependent estimates of $(\mathbf{J}_{ii}, \sigma_i^2, b_i)$. (iii) *Metric-aware β .* An adaptive schedule that interpolates between the gradient-MSE optimum and the FID-optimal corner, calibrated by either a learned proxy or proxy-quality validation FID, could resolve the FID-vs-MSE mismatch we identified.

References

- [1] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 6840–6851. Curran Associates, Inc., 2020.
- [2] Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations*, 2021.
- [3] Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. In *Advances in Neural Information Processing Systems*, volume 35, 2022.
- [4] Prafulla Dhariwal and Alexander Nichol. Diffusion models beat GANs on image synthesis. In *Advances in Neural Information Processing Systems*, volume 34, 2021.
- [5] Patrick Esser, Sumith Kulal, Andreas Blattmann, Rahim Entezari, Jonas Müller, Harry Saini, Yam Levi, Dominik Lorenz, Naveen Rafi, Tim Shafir, et al. Scaling rectified flow transformers for high-resolution image synthesis. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, 2024.
- [6] Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David K. Duvenaud. Neural ordinary differential equations. In *Advances in Neural Information Processing Systems*, volume 31, 2018.

- [7] Danilo Rezende and Shakir Mohamed. Variational inference with normalizing flows. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pages 1530–1538, 2015.
- [8] Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density estimation using Real-NVP. In *International Conference on Learning Representations*, 2017.
- [9] Will Grathwohl, Ricky T. Q. Chen, Jesse Bettencourt, Ilya Sutskever, and David Duvenaud. FFJORD: Free-form continuous dynamics for scalable reversible generative models. In *International Conference on Learning Representations*, 2019.
- [10] Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling, 2023.
- [11] Alexander Tong, Kilian Fatras, Nikolay Malkin, Guillaume Huguet, Yanlei Zhang, Jarrod Rector-Brooks, Guy Wolf, and Yoshua Bengio. Improving and generalizing flow-based generative models with minibatch optimal transport, 2024.
- [12] Michael S. Albergo and Eric Vanden-Eijnden. Building normalizing flows with stochastic interpolants, 2023.
- [13] Michael Albergo, Nicholas M. Boffi, and Eric Vanden-Eijnden. Stochastic interpolants: A unifying framework for flows and diffusions. *Journal of Machine Learning Research*, 26(209):1–80, 2025.
- [14] Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- [15] Tim Salimans and Jonathan Ho. Progressive distillation for fast sampling of diffusion models, 2022.
- [16] Tianwei Yin, Michaël Gharbi, Richard Zhang, Eli Shechtman, Frédo Durand, William T. Freeman, and Taesung Park. One-step diffusion with distribution matching distillation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 6613–6623, June 2024.
- [17] Tianwei Yin, Michaël Gharbi, Taesung Park, Richard Zhang, Eli Shechtman, Frédo Durand, and William T. Freeman. Improved distribution matching distillation for fast image synthesis. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 47455–47487. Curran Associates, Inc., 2024.
- [18] Yang Song, Prafulla Dhariwal, Mark Chen, and Ilya Sutskever. Consistency models. In *Proceedings of the 40th International Conference on Machine Learning, ICML’23*. JMLR.org, 2023.
- [19] Cheng Lu and Yang Song. Simplifying, stabilizing and scaling continuous-time consistency models, 2025.
- [20] Nicholas M. Boffi, Michael S. Albergo, and Eric Vanden-Eijnden. Flow map matching with stochastic interpolants: A mathematical framework for consistency models, 2025.
- [21] Zhengyang Geng, Mingyang Deng, Xingjian Bai, J. Zico Kolter, and Kaiming He. Mean flows for one-step generative modeling, 2025.
- [22] Huijie Zhang, Aliaksandr Siarohin, Willi Menapace, Michael Vasilkovsky, Sergey Tulyakov, Qing Qu, and Ivan Skorokhodov. Alphaflow: Understanding and improving meanflow models, 2025.
- [23] Zhengyang Geng, Yiyang Lu, Zongze Wu, Eli Shechtman, J. Zico Kolter, and Kaiming He. Improved mean flows: On the challenges of fastforward generative models, 2025.
- [24] Xinxi Zhang, Shiwei Tan, Quang Nguyen, Quan Dao, Ligong Han, Xiaoxiao He, Tunyu Zhang, Chengzhi Mao, Dimitris Metaxas, and Vladimir Pavlovic. Overcoming the curvature bottleneck in meanflow, 2026.
- [25] Linqi Zhou, Mathias Parger, Ayaan Haque, and Jiaming Song. Terminal velocity matching, 2026.
- [26] Zhiqi Li, Yuchen Sun, Greg Turk, and Bo Zhu. Functional mean flow in hilbert space, 2025.
- [27] Paul Glasserman. *Monte Carlo methods in financial engineering*, volume 53. Springer New York, NY, 2003.
- [28] Haochen You, Baojing Liu, and Hongyang He. Modular meanflow: Towards stable and scalable one-step generative modeling, 2025.

- [29] Jin-Young Kim, Hyojun Go, Lea Bogensperger, Julius Erbach, Nikolai Kalischek, Federico Tombari, Konrad Schindler, and Dominik Narnhofer. Understanding, accelerating, and improving meanflow training, 2025.
- [30] Kyungmin Lee, Sihyun Yu, and Jinwoo Shin. Decoupled meanflow: Turning flow models into flow maps for accelerated sampling, 2025.
- [31] Dongyeop Woo, Marta Skreta, Seonghyun Park, Kirill Neklyudov, and Sungsoo Ahn. Riemannian meanflow, 2026.
- [32] Elaheh Akbari, Shansita Sharma, Ping He, Ahmadreza Moradipari, Kyungtae Han, Hamed Pirsiavash, Yikun Bai, and Soheil Kolouri. OT-MeanFlow3D: Bridging optimal transport and meanflow for efficient 3D point cloud generation, 2025.
- [33] Donglin Yang, Yongxing Zhang, Xin Yu, Liang Hou, Xin Tao, Pengfei Wan, Xiaojuan Qi, and Renjie Liao. Stable velocity: A variance perspective on flow matching, 2026.
- [34] Chika Maduabuchi and Jindong Wang. Temporal pair consistency for variance-reduced flow matching, 2026.
- [35] Shadab Ahamed, Eshed Gal, Simon Ghyselincks, Md Shahriar Rahim Siddiqui, Moshe Eliasof, and Eldad Haber. Preconditioned score and flow matching, 2026.
- [36] Quentin Bertrand, Anne Gagneux, Mathurin Massias, and Rémi Emonet. On the closed-form of flow matching: Generalization does not arise from target stochasticity, 2025.
- [37] Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow, 2022.
- [38] Sangyun Lee, Zinan Lin, and Giulia Fanti. Improving the training of rectified flows. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 63082–63109. Curran Associates, Inc., 2024.
- [39] Peter W. Glynn and Roberto Szechtman. Some new perspectives on the method of control variates. *Monte Carlo and Quasi-Monte Carlo Methods 2000*, pages 27–49, 2002.
- [40] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, et al. Human-level control through deep reinforcement learning. *Nature*, 518:529–533, 2015.
- [41] Jean-Bastien Grill, Florian Strub, Florent Altché, et al. Bootstrap your own latent: A new approach to self-supervised learning. In *NeurIPS*, 2020.
- [42] William Peebles and Saining Xie. Scalable diffusion models with transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4195–4205, 2023.
- [43] Kevin Frans, Danijar Hafner, Sergey Levine, and Pieter Abbeel. One step diffusion via shortcut models, 2025.
- [44] Cédric Villani et al. *Optimal Transport: Old and New*, volume 338. Springer Berlin, Heidelberg, 2009.
- [45] Boris T. Polyak and Anatoli B. Juditsky. Acceleration of stochastic approximation by averaging. volume 30, pages 838–855, 1992.
- [46] Antti Tarvainen and Harri Valpola. Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised learning results. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The claims in the abstract and introduction highlight the three primary contributions explicitly, including the control-variate diagnosis for the MeanFlow’s training pathology and the closed-form optimal coefficient in section 3, and the numerical results from our empirical study in section 5.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We explicitly discuss our limitations in a paragraph of section 6.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: All theorems and propositions we establish in section 3 are numbered, cross-referenced, and come with a clear description of the assumptions. We provide complete and extensive proofs in the supplementary material.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All the primary practical information to run the two-dimensional experiments and DiT on ImageNet are clearly stated in the main text, with further details given in the supplementary material.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in

some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We will release the JAX/Flax implementation for our primary experiments upon the acceptance of this paper.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: We have an experiment setup paragraph in section 5 elaborating all the details for the experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: For toy and DGMM experiments, every reported number aggregates three random seeds $\{42, 0, 1\}$ with the standard error of the mean reported alongside (Tables 3 and 4 and figs. 2 and 3). The DiT-B/4 / ImageNet-256 experiments are single-seed (acknowledged in the Limitations paragraph of Section 6); the matched-step gap analysis in Appendix F reports asymptotic offsets averaged over 11 consecutive matched-step windows (± 0.11 FID standard deviation), which provides a within-run measure of consistency rather than across-seed variance.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: We report all the information about the resources to run the experiments in the supplementary material.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The study is theoretical and methodological and compliant with the NeurIPS code of ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The study in this paper comes with a theoretical nature. There is no direct broader impact to declare at this moment.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: This paper does not pose any of such risks.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: We cited all external assets in the references, including the ImageNet dataset, the stable Diffusion VAE for inference, the DiT architecture, the original MeanFlow algorithm, and concurrent works.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not release any new asset at this moment.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: Neither does this paper involve any crowdsourcing nor research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

810 Answer: [N/A]

811 Justification: Neither does this paper involve any crowdsourcing nor research with human

812 subjects.

813 Guidelines:

- 814 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
- 815 with human subjects.
- 816 • Depending on the country in which research is conducted, IRB approval (or equiva-
- 817 lent) may be required for any human subjects research. If you obtained IRB approval,
- 818 you should clearly state this in the paper.
- 819 • We recognize that the procedures for this may vary significantly between institutions
- 820 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
- 821 guidelines for their institution.
- 822 • For initial submissions, do not include any information that would break anonymity
- 823 (if applicable), such as the institution conducting the review.

824 **16. Declaration of LLM usage**

825 Question: Does the paper describe the usage of LLMs if it is an important, original, or

826 non-standard component of the core methods in this research? Note that if the LLM is used

827 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

828 scientific rigor, or originality of the research, declaration is not required.

829 Answer: [N/A]

830 Justification: This paper does not involve LLMs as a primary component.

831 Guidelines:

- 832 • The answer [N/A] means that the core method development in this research does not
- 833 involve LLMs as any important, original, or non-standard components.
- 834 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
- 835 not be described.

836	Supplementary Material	
837	A Notations	22
838	B Toy and DGMM Benchmark Details	22
839	C Proof of Lemma	23
840	D Implementation Details	24
841	E Proofs of Theorems and Propositions	25
842	E.1 Proof of Theorem 2 (Jacobian Variance Amplification)	25
843	E.2 Proof of Theorem 3 (Semi-Gradient Gap)	26
844	E.3 Proof of Proposition 2 (Gradient under the $\beta = 1$ corner loss)	26
845	E.4 Proof of Proposition 3 (Bias-Variance Tradeoff)	27
846	E.5 Proof of Proposition 4 (Target/Tangent Bias Asymmetry)	27
847	E.6 Proof of Theorem 4 (Optimal Control-Variate Coefficient)	28
848	F DiT Latent-Space Training on ImageNet	29
849	F.1 DiT-B/4, ImageNet- 256×256	29

850 A Notations

Table 2 lists all the notations we use in this work.

Table 2: Table of notations.

Notation	Description
<i>States and velocities</i>	
$\mathbf{x}, \mathbf{x}_t$	Random data sample / state at interpolant time t
$\mathbf{x}, \mathbf{x}_t$	Realized state / state at time t : $\mathbf{x}_t = (1 - t)\mathbf{x}_0 + t\mathbf{x}_1$
$\mathbf{v}(\mathbf{x}, t)$	Marginal velocity field
$\mathbf{v}_{\text{cond}}$	Conditional velocity: $\mathbf{v}_{\text{cond}} = \mathbf{x}_1 - \mathbf{x}_0$
$\mathbf{v}'$	Conditional velocity fluctuation: $\mathbf{v}' = \mathbf{v}_{\text{cond}} - \mathbf{v}(\mathbf{x}_t, t)$, $\mathbb{E}_{\mathbf{x}_0 \mathbf{x}_t}[\mathbf{v}'] = \mathbf{0}$
$\hat{\mathbf{v}}$	Deterministic proxy for $\mathbf{v}$ (e.g., EMA velocity at boundary $r = t$)
<i>Networks</i>	
$\mathbf{u}_{\theta}(\mathbf{x}, r, t)$	Learned two-parameter average velocity field (parameters θ)
$\mathbf{u}_{\hat{\theta}}$	EMA copy of $\mathbf{u}_{\theta}$, parameters $\hat{\theta}$
<i>Spatial-Jacobian quantities</i>	
$\partial_{\mathbf{z}}\mathbf{u}_{\theta}$	Spatial Jacobian of $\mathbf{u}_{\theta}$ w.r.t. its first argument
$\mathbf{J}$	Jacobi factor: $\mathbf{J} = (t - r)\partial_{\mathbf{z}}\mathbf{u}_{\theta} - \mathbf{I}_d \in \mathbb{R}^{d \times d}$
$\mathbf{A}$	Shorthand $\mathbf{A} \triangleq \mathbf{J} + \mathbf{I}_d = (t - r)\partial_{\mathbf{z}}\mathbf{u}_{\theta}$
κ	Scalar-isotropic surrogate for $\mathbf{J}$: $\mathbf{J} \approx \kappa\mathbf{I}_d$
<i>Bias-variance ingredients</i>	
$\Sigma_{\mathbf{v}'}$	Conditional velocity covariance: $\text{Cov}_{\mathbf{x}_0 \mathbf{x}_t}[\mathbf{v}']$
$\sigma^2 d$	Trace of $\Sigma_{\mathbf{v}'}$; in scalar-isotropic case $\Sigma_{\mathbf{v}'} = \sigma^2 \mathbf{I}_d$
$\mathbf{b}$	Proxy bias: $\mathbf{b} \triangleq \hat{\mathbf{v}} - \mathbf{v}(\mathbf{x}_t, t)$
$\ \mathbf{b}\ ^2$	EMA-tracking proxy: $\mathbb{E}_{\mathbf{x}_t}[\ \mathbf{u}_{\theta}(\mathbf{x}_t, t, t) - \mathbf{u}_{\hat{\theta}}(\mathbf{x}_t, t, t)\ ^2]$
<i>Gradients</i>	
$\nabla_{\theta}\ell_{\text{MF}}$	Per-step gradient of the MeanFlow loss
$\mathbf{g}$	Parameter Jacobian of the network output: $\mathbf{g} \triangleq \nabla_{\theta}\mathbf{u}_{\theta} \in \mathbb{R}^{d \times p}$
$\mathbf{G}_{\theta}$	Parameter-space Gram matrix: $\mathbf{G}_{\theta} \triangleq \mathbf{g}\mathbf{g}^{\top} \in \mathbb{R}^{d \times d}$
<i>Tangent-mixing coefficient</i>	
$\beta \in [0, 1]$	Tangent-mixing coefficient equation 6
β^*	Closed-form scalar-isotropic minimizer of $M(\beta)$ in Theorem 4
β_{matrix}^*	Matrix-form minimizer (Appendix E, equation equation 25)
$M(\beta)$	Conditional MSE of the per-sample gradient equation 7 at coefficient β
<i>Loss components</i>	
$\mathcal{L}_{\text{MF}}$	Vanilla MeanFlow loss equation 4
$\mathcal{L}_{\text{MF}}^{\text{EMA}}$	EMA-tangent variant of the MeanFlow loss equation 9
$\mathcal{L}_{\text{FM}}$	Flow-matching anchor loss equation 10
$\mathcal{L}_{\beta=1}$	$\beta = 1$ corner training loss equation 11
<i>Operators</i>	
$\text{sg}[\cdot]$	Stop-gradient operator
$\text{JVP}(\mathbf{u}, \mathbf{x}, \mathbf{v})$	Jacobian-vector product of $\mathbf{u}$ at $\mathbf{x}$ in tangent direction $\mathbf{v}$
$\text{Tr}(\cdot), \text{Cov}[\cdot], \text{Var}[\cdot]$	Trace, covariance, variance

851

852 B Toy and DGMM Benchmark Details

853 **Toy 2-D benchmark.** Table 3 reports the matched-method comparison of vanilla MeanFlow
854 ($\beta = 0$) against the full $\beta = 1$ corner recipe (EMA tangent, FM anchor with $(\delta_{\min}, \delta_{\max}, \lambda) =$
855 $(10^{-4}, 10^{-2}, 0.5)$, no trace weight, no adaptive weighting) on six 2-D datasets, averaged over three
856 seeds at 200k optimization steps. The $\beta = 1$ recipe reduces SW_1 on five of six datasets, with the

Table 3: Sliced Wasserstein distance (lower is better) on the 2-D toy benchmark, averaged over three seeds $\{42, 0, 1\}$ at 200k steps. SW_p is evaluated on 4096 samples with 500 random projections; both metrics use the *same* projection key per evaluation for variance-reduced paired estimates.

Dataset	Metric	MeanFlow ($\beta=0$)	$\beta=1$
checkerboard	SW_1	0.175 ± 0.036	0.134 ± 0.029
	SW_2	0.234 ± 0.042	0.175 ± 0.046
eight_gaussians	SW_1	0.098 ± 0.016	0.085 ± 0.010
	SW_2	0.151 ± 0.022	0.146 ± 0.023
two_moons	SW_1	0.093 ± 0.028	0.072 ± 0.019
	SW_2	0.177 ± 0.068	0.093 ± 0.025
swiss_roll	SW_1	0.098 ± 0.058	0.045 ± 0.003
	SW_2	0.177 ± 0.151	0.060 ± 0.007
two_spirals	SW_1	0.034 ± 0.001	0.037 ± 0.001
	SW_2	0.044 ± 0.001	0.046 ± 0.001
pinwheel	SW_1	0.094 ± 0.011	0.059 ± 0.005
	SW_2	0.128 ± 0.014	0.076 ± 0.005

Table 4: DGMM scaling: SW_1 at convergence as the dimension d grows (mean $\pm$ standard error over three seeds, 200k steps). Methods are within one standard error at every d .

d	2	4	8	16	32	64
MeanFlow ($\beta=0$)	0.152 ± 0.018	0.083 ± 0.009	0.062 ± 0.003	0.044 ± 0.002	0.034 ± 0.001	0.028 ± 0.003
$\beta=1$ corner	0.147 ± 0.012	0.083 ± 0.009	0.061 ± 0.002	0.043 ± 0.002	0.036 ± 0.002	0.028 ± 0.001

largest reduction (54%) on `swiss_roll`. `two_spirals` is a known low-curvature outlier where vanilla MeanFlow already attains the lowest absolute SW_1 ; the empirical $\beta^* = 0$ on the β -sweep agrees.

High-dimensional scaling on DGMM (matched-method comparison). The matched-method comparison of $\beta = 0$ and $\beta = 1$ in Table 4 (a separate 2-method, 3-seed run independent of the 11-point β -sweep of Section 5.3) was run before the full β -sweep. The two configurations are statistically indistinguishable at every d , consistent with the full β -sweep finding that the bias-variance landscape on DGMM is empirically near-flat at low d (the EMA proxy’s $\|\mathbf{b}\|^2$ is near-zero on approximately isotropic data, so the corner-versus-interior gap is within-noise).

C Proof of Lemma

Lemma 1. *Given conditional vector fields $\mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0)$ that generate probability paths $p(\mathbf{x}, t \mid \mathbf{x}_0)$, for any data distribution $\mathbf{x}_0 \sim p(\mathbf{x}_0)$, define the marginal velocity field as*

$$\mathbf{v}(\mathbf{x}, t) \triangleq \mathbb{E}_{\mathbf{x}_0 \sim p(\mathbf{x}_0 \mid \mathbf{x}_t = \mathbf{x})}[\mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0)].$$

Then $\mathbf{v}(\mathbf{x}, t)$ generates the marginal probability path $p(\mathbf{x}, t) = \int p(\mathbf{x}, t \mid \mathbf{x}_0) p(\mathbf{x}_0) d\mathbf{x}_0$ via the continuity equation.

Proof. We follow the standard flow-matching construction [10, 11]: the marginal velocity field is defined by the displayed conditional expectation, and it remains to verify that this definition makes $\mathbf{v}(\mathbf{x}, t)$ generate $p(\mathbf{x}, t)$ via the continuity equation. By the law of total probability and the definition of the conditional velocity field,

$$p(\mathbf{x}, t) \mathbf{v}(\mathbf{x}, t) = \int p(\mathbf{x}, t \mid \mathbf{x}_0) \mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0) p(\mathbf{x}_0) d\mathbf{x}_0. \quad (8)$$

Each conditional pair $(\mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0), p(\mathbf{x}, t \mid \mathbf{x}_0))$ satisfies the continuity equation $\partial_t p(\mathbf{x}, t \mid \mathbf{x}_0) + \text{div}(p(\mathbf{x}, t \mid \mathbf{x}_0) \mathbf{v}(\mathbf{x}, t \mid \mathbf{x}_0)) = 0$ by assumption [44]. Integrating against $p(\mathbf{x}_0)$ and exchanging

Algorithm 1 Training the $\beta = 1$ corner instantiation (EMA tangent + FM anchor)

Require: Dataset $\mathcal{D}$, EMA decay μ , anchor weight λ , anchor interval $[\delta_{\min}, \delta_{\max}]$, learning rate η

Initialize $\bar{\theta} \leftarrow \theta$

for $k = 1, 2, \dots$ **do**

Sample $(x_0, x_1) \sim \mathcal{D} \times \mathcal{N}(\mathbf{0}, \mathbf{I}_d)$, $t \sim \mathcal{U}[0, 1]$, $r \sim \mathcal{U}[0, t]$, $\delta \sim \mathcal{U}[\delta_{\min}, \delta_{\max}]$

$x_t \leftarrow (1 - t)x_0 + tx_1$, $v_{\text{cond}} \leftarrow x_1 - x_0$

$v_{\text{tang}} \leftarrow \text{sg}[\mathbf{u}_{\bar{\theta}}(x_t, t, t)]$ $\triangleright$ EMA tangent (one extra forward pass)

$V_{\theta} \leftarrow \mathbf{u}_{\theta}(x_t, r, t) + (t - r) \text{sg}[\text{JVP}(\mathbf{u}_{\theta}, (x_t, r, t), (v_{\text{tang}}, 0, 1))]$

$\ell_{\text{MF}} \leftarrow \|V_{\theta} - v_{\text{cond}}\|_2^2$

$\ell_{\text{FM}} \leftarrow \|\mathbf{u}_{\theta}(x_t, t - \delta, t) - v_{\text{cond}}\|_2^2$ $\triangleright$ FM anchor (one extra forward pass)

$\theta \leftarrow \theta - \eta \nabla_{\theta}(\ell_{\text{MF}} + \lambda \ell_{\text{FM}})$

$\bar{\theta} \leftarrow \mu \bar{\theta} + (1 - \mu)\theta$ $\triangleright$ EMA update

end for

return θ $\triangleright$ Generate: $\hat{x}_0 = x_1 - \mathbf{u}_{\theta}(x_1, 0, 1)$

integration with ∂_t and div yields $\partial_t p(x, t) + \text{div}(p(x, t)v(x, t)) = 0$, so the constructed $v(x, t)$ generates $p(x, t)$. Dividing the displayed equation by $p(x, t)$ and applying Bayes' rule $p(x_0 | x_t = x) = p(x, t | x_0)p(x_0)/p(x, t)$ recovers the conditional-expectation form. We treat this as the operating definition of the marginal velocity used throughout; uniqueness up to divergence-free vector fields, which we do not require, is a separate question we do not address. $\square$

D Implementation Details

To validate the framework empirically (Section 5), we instantiate the $\beta = 1$ corner of Theorem 4 as a backbone-preserving training recipe. Two design choices realize the corner at near-zero compute overhead.

EMA proxy. We take $\hat{v} = \mathbf{u}_{\bar{\theta}}(x_t, t, t)$ where $\bar{\theta}$ is a Polyak-averaged copy of θ [45, 46], $\bar{\theta} \leftarrow \mu \bar{\theta} + (1 - \mu)\theta$ with $\mu \in [0, 1]$. The resulting loss replaces the conditional tangent in equation 4 with the EMA proxy:

$$\mathcal{L}_{\text{MF}}^{\text{EMA}}(\theta) = \mathbb{E}_{r, t, x_0, x_1} \left[\left\| \mathbf{u}_{\theta} + (t - r) \text{sg}[\text{JVP}(\mathbf{u}_{\theta}, (z, r, t), (\mathbf{u}_{\bar{\theta}, t}, 0, 1))] - v_{\text{cond}} \right\|_2^2 \right], \quad (9)$$

with $\mathbf{u}_{\bar{\theta}, t} \triangleq \mathbf{u}(x_t, t, t; \bar{\theta})$. Crucially, the regression target remains v_{cond} : the tangent and target play distinct statistical roles per Section 3.1, and only the tangent needs to be replaced.

Flow-matching anchor. The proxy bias $\mathbf{b} = \mathbf{u}_{\bar{\theta}, t} - v$ must be controlled or it propagates through the $\beta(\mathbf{J} + \mathbf{I})\mathbf{b}$ bias term. We supervise the boundary condition $\mathbf{u}(x_t, t, t) = v(x_t, t)$ directly with a small flow-matching loss

$$\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{t, x_0, x_1} \left[\left\| \mathbf{u}_{\theta}(x_t, t - \delta, t) - v_{\text{cond}} \right\|_2^2 \right], \quad (10)$$

where $\delta \sim \mathcal{U}[\delta_{\min}, \delta_{\max}]$ with $\delta_{\max} \ll 1$. We chose to evaluate the anchor at a small *interior* time offset rather than exactly at the boundary $\delta = 0$ for two reasons: (i) the same $\mathbf{u}(x_t, t - \delta, t)$ is what the EMA tangent eventually converges to, so the anchor and the JVP target share the same network-evaluation pattern, and (ii) we found training to be slightly more stable when the anchor probes a thin neighborhood of the boundary rather than a single $r = t$ slice. For small δ , $\mathbf{u}(x_t, t - \delta, t)$ approximates the instantaneous velocity, and equation 1 guarantees $\mathbb{E}[v_{\text{cond}} | x_t] = v(x_t, t)$, so $\mathcal{L}_{\text{FM}}$ drives $\mathbf{u}_{\theta}(x_t, t, t) \rightarrow v(x_t, t)$, and via EMA, $\mathbf{u}_{\bar{\theta}}(x_t, t, t) \rightarrow v$ and hence $\mathbf{b} \rightarrow \mathbf{0}$. The full $\beta = 1$ corner loss is

$$\mathcal{L}_{\beta=1}(\theta) = \mathcal{L}_{\text{MF}}^{\text{EMA}}(\theta) + \lambda \mathcal{L}_{\text{FM}}(\theta), \quad (11)$$

with $\lambda > 0$. The training procedure is Algorithm 1. Anchor hyperparameters used in our experiments: for 2-D toy and DGMM runs, $(\delta_{\min}, \delta_{\max}, \lambda) = (10^{-4}, 10^{-2}, 0.5)$; for the DiT-B/4 / ImageNet-256 $\beta = 1$ run, $(\delta_{\min}, \delta_{\max}, \lambda) = (0, 10^{-3}, 0.1)$.

905 **Optimality.** The combination of EMA proxy plus FM anchor approximates the $\beta \rightarrow 1$ limit of The-
 906 orem 4 at the cost of one additional JVP per step. Propositions 2 and 3 (proofs in Appendix E)
 907 characterize the resulting gradient and bias structure:

908 **Proposition 2** (Gradient under the $\beta = 1$ corner loss). *The gradient of $\mathcal{L}_{MF}^{EMA}$ takes the form*

$$\nabla_{\theta} \mathcal{L}_{MF}^{EMA} = 2 \mathbb{E}[g \tilde{r}^{EMA}],$$

909 where $(\nabla_{\theta} \mathbf{u}_{\theta}) \in \mathbb{R}^{d \times p}$ is the parameter Jacobian, $\tilde{r}^{EMA} \triangleq V_{\theta}^{EMA} - \mathbf{v}(\mathbf{x}_t, t) \in \mathbb{R}^d$, and no $\mathbf{J} \mathbf{v}'$ noise
 910 term appears.

911 **Proposition 3** (Bias-Variance Tradeoff). *The EMA mean-field residual decomposes as*

$$\tilde{r}^{EMA} = \mathbf{r}_{\theta} + (t-r) \partial_z \mathbf{u}_{\theta}(\mathbf{u}_{\hat{\theta}}(\mathbf{x}_t, t, t) - \mathbf{v}(\mathbf{x}_t, t)),$$

912 where $\mathbf{r}_{\theta}$ is the true MeanFlow residual. The FM anchor drives $\mathbf{u}_{\hat{\theta}}(\mathbf{x}_t, t, t) \rightarrow \mathbf{v}(\mathbf{x}_t, t)$, giving
 913 $\tilde{r}^{EMA} \rightarrow \mathbf{0}$.

914 **Why the target stays unbiased.** Propositions 2 and 3 treat the proxy $\hat{\mathbf{v}}$ as a tangent-only re-
 915 placement. A natural alternative is to substitute $\hat{\mathbf{v}}$ for the regression target $\mathbf{v}_{\text{cond}}$ as well, yielding a
 916 “fully self-consistent” loss in the spirit of self-distillation. The next proposition shows that the two
 917 replacements are not interchangeable: only the tangent admits a gap-attenuated bias.

918 **Proposition 4** (Target/Tangent Bias Asymmetry). *Let $\hat{\mathbf{v}} = \mathbf{u}_{\hat{\theta}}(\mathbf{x}_t, t, t)$ be a deterministic proxy
 919 with bias $\mathbf{b} \triangleq \hat{\mathbf{v}} - \mathbf{v}(\mathbf{x}_t, t)$. Consider two ways of substituting $\hat{\mathbf{v}}$ into equation 4:*

920 (i) Tangent replacement (the $\beta = 1$ corner): *keep the target $\mathbf{v}_{\text{cond}}$ but replace the JVP tangent.*
 921 *At the loss minimum,*

$$\mathbf{u}^*(\mathbf{x}_t, r, t) - \mathbf{v}(\mathbf{x}_t, t) = -\frac{1}{2}(t-r) \partial_z \mathbf{u}^* \mathbf{b} + \mathcal{O}((t-r) \|\mathbf{b}\|^2),$$

922 *so the stationary error is $\mathcal{O}((t-r) \|\partial_z \mathbf{u}\| \|\mathbf{b}\|)$ and vanishes at $r \rightarrow t$.*

923 (ii) Target replacement: *substitute $\hat{\mathbf{v}}$ for $\mathbf{v}_{\text{cond}}$ in the regression target. At $r = t$ the loss minimum*
 924 *satisfies*

$$\mathbf{u}^*(\mathbf{x}_t, t, t) - \mathbf{v}(\mathbf{x}_t, t) = \mathbf{b},$$

925 *i.e., the stationary error equals the proxy bias exactly, uniformly in t and independent of*
 926 *the gap $(t-r)$.*

927 The proof is given in Appendix E. The asymmetry has three immediate consequences. *First*, the
 928 $\beta = 1$ corner’s design choice (replace only the tangent, keep $\mathbf{v}_{\text{cond}}$ as target) is not arbitrary: tangent-
 929 replacement bias is multiplicatively damped by the gap $(t-r)$, while target-replacement bias is not.
 930 *Second*, the FM anchor in equation 10 suffices to control the residual tangent-replacement bias
 931 because it pins $\mathbf{u}_{\theta}(\mathbf{x}_t, t, t) \rightarrow \mathbf{v}$ at the boundary $r = t$, where any target-replacement scheme would
 932 require an analogous anchor at *every* (r, t) . *Third*, the proposition locates a failure mode in self-
 933 distillation-style MeanFlow variants that use a deterministic proxy for both roles: the loss admits
 934 the trivial fixed point $\mathbf{u}^* \equiv \hat{\mathbf{v}}$, which is preserved under proxy updates and only escaped by external
 935 supervision pulling $\hat{\mathbf{v}}$ back toward $\mathbf{v}$ at every t . Concurrent methods that adopt this pattern indeed
 936 report sensitivity to the proxy-update schedule and the boundary regularizer [23, 24].

937 Proposition 2 confirms that the $\mathbf{J} \mathbf{v}'$ amplification of Theorem 2 is eliminated; Proposition 3 confirms
 938 that the residual EMA bias is controlled by the FM anchor, recovering the asymptotic optimum
 939 of Theorem 4; Proposition 4 explains why the FM anchor at the boundary $r = t$ alone is sufficient.

940 E Proofs of Theorems and Propositions

941 E.1 Proof of Theorem 2 (Jacobian Variance Amplification)

942 **Theorem 5** (Jacobian Variance Amplification). *The Jacobi factor amplifies the gradient variance*
 943 *of the original MeanFlow loss:*

$$\text{Var}[\nabla_{\theta} \ell_{MF}(\theta) \mid \mathbf{x}_t] \propto \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^{\top}).$$

944 *Proof.* Under the stop-gradient operator, the compound prediction V_θ^{sg} depends on θ only through
 945 $\mathbf{u}_\theta(\mathbf{x}_t, r, t)$. Let $\nabla_\theta \mathbf{u}_\theta \in \mathbb{R}^{d \times p}$ denote the parameter Jacobian of the network output. The per-step
 946 gradient is

$$\nabla_\theta \ell_{\text{MF}} = 2(\nabla_\theta \mathbf{u}_\theta)^\top (\mathbf{r}_\theta^{\text{sg}} + \mathbf{J} \mathbf{v}'), \quad (12)$$

947 where $\mathbf{r}_\theta^{\text{sg}}$ is the mean-field residual and $\mathbf{J} = (t - r) \partial_z \mathbf{u}_\theta - \mathbf{I}_d$ is the Jacobi factor. All three
 948 quantities ($\nabla_\theta \mathbf{u}_\theta$, $\mathbf{r}_\theta^{\text{sg}}$, and $\mathbf{J}$) are deterministic conditioned on $\mathbf{x}_t$.

949 Since $\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t}[\mathbf{v}'] = \mathbf{0}$ by equation 1, the conditional mean gradient is:

$$\mathbb{E}[\nabla_\theta \ell | \mathbf{x}_t] = 2(\nabla_\theta \mathbf{u}_\theta)^\top \mathbf{r}_\theta^{\text{sg}}. \quad (13)$$

950 The centered gradient is $\nabla_\theta \ell - \mathbb{E}[\nabla_\theta \ell | \mathbf{x}_t] = 2(\nabla_\theta \mathbf{u}_\theta)^\top \mathbf{J} \mathbf{v}'$, so the conditional variance is:

$$\begin{aligned} \text{Var}[\nabla_\theta \ell | \mathbf{x}_t] &= \mathbb{E}[\|2(\nabla_\theta \mathbf{u}_\theta)^\top \mathbf{J} \mathbf{v}'\|^2 | \mathbf{x}_t] \\ &= 4 \mathbb{E} \left[(\mathbf{v}')^\top \mathbf{J}^\top \underbrace{(\nabla_\theta \mathbf{u}_\theta)(\nabla_\theta \mathbf{u}_\theta)^\top}_{\mathbf{G}_\theta \succeq \mathbf{0}} \mathbf{J} \mathbf{v}' \mid \mathbf{x}_t \right] \\ &= 4 \text{Tr}(\mathbf{G}_\theta \mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top) \propto \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top), \end{aligned} \quad (14)$$

951 where $\mathbf{G}_\theta = (\nabla_\theta \mathbf{u}_\theta)(\nabla_\theta \mathbf{u}_\theta)^\top \in \mathbb{R}^{d \times d}$ is the parameter-space Gram matrix and the last step uses
 952 $\mathbf{G}_\theta \succeq \mathbf{0}$ to establish proportionality via the cyclic property of the trace. $\square$

953 E.2 Proof of Theorem 3 (Semi-Gradient Gap)

954 **Theorem 6** (Semi-Gradient Gap). *The gradients of the MeanFlow loss with and without the stop-*
 955 *gradient operator differ by*

$$\underbrace{2(t-r) \mathbb{E}_{r,t,\mathbf{x}_0,\mathbf{x}_1} [(\nabla_\theta(\partial_z \mathbf{u}_\theta \cdot \mathbf{v}(\mathbf{x}_t, t) + \partial_t \mathbf{u}_\theta))^\top \mathbf{r}_\theta]}_{\text{mean-field gradient difference}} + \underbrace{\mathbb{E}_{r,t,\mathbf{x}_0,\mathbf{x}_1} [\nabla_\theta \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top)]}_{\text{variance-driven gradient difference}}.$$

956 *Proof.* The MeanFlow loss conditioned on $\mathbf{x}_t$ decomposes as (cf. equation 5):

$$\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t}[\ell_{\text{MF}}] = \|\mathbf{r}_\theta^{\text{sg}}\|^2 + \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top). \quad (15)$$

957 **Without stop-gradient.** Both terms depend on θ (through $\mathbf{u}_\theta$ and $\partial_z \mathbf{u}_\theta$), so the full gradient is:

$$\nabla_\theta \mathbb{E}[\ell] = \nabla_\theta \|\mathbf{r}_\theta\|^2 + \nabla_\theta \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top). \quad (16)$$

958 **With stop-gradient.** The JVP terms in $\mathbf{r}_\theta^{\text{sg}}$ and $\mathbf{J}$ are treated as constants, so:

$$\nabla_\theta^{\text{sg}} \|\mathbf{r}_\theta^{\text{sg}}\|^2 = 2(\nabla_\theta \mathbf{u}_\theta)^\top \mathbf{r}_\theta^{\text{sg}}, \quad \nabla_\theta^{\text{sg}} \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top) = \mathbf{0}. \quad (17)$$

959 The gap between the full and stop-gradient gradients is therefore:

$$\begin{aligned} \nabla_\theta \mathbb{E}[\ell] - \nabla_\theta^{\text{sg}} \mathbb{E}[\ell] &= \underbrace{\nabla_\theta \|\mathbf{r}_\theta\|^2 - 2(\nabla_\theta \mathbf{u}_\theta)^\top \mathbf{r}_\theta^{\text{sg}}}_{\text{mean-field gradient difference}} + \underbrace{\nabla_\theta \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top)}_{\text{variance-driven difference}}. \end{aligned} \quad (18)$$

960 Expanding the first term: $\nabla_\theta \|\mathbf{r}_\theta\|^2 = 2(\nabla_\theta \mathbf{r}_\theta)^\top \mathbf{r}_\theta$, and $\nabla_\theta \mathbf{r}_\theta$ includes the derivative of the JVP
 961 terms $(t - r)(\partial_z \mathbf{u}_\theta \cdot \mathbf{v} + \partial_t \mathbf{u}_\theta)$ w.r.t. θ , yielding the second-order mean-field difference $2(t -$
 962 $r) \mathbb{E}[(\nabla_\theta(\partial_z \mathbf{u}_\theta \cdot \mathbf{v} + \partial_t \mathbf{u}_\theta))^\top \mathbf{r}_\theta]$. At convergence $\mathbf{r}_\theta \rightarrow \mathbf{0}$, this term vanishes, and the gap is
 963 dominated by the variance-driven term $\nabla_\theta \text{Tr}(\mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top)$. $\square$

964 E.3 Proof of Proposition 2 (Gradient under the $\beta = 1$ corner loss)

965 **Proposition 7** (Gradient under the $\beta = 1$ corner loss). *The gradient of $\mathcal{L}_{\text{MF}}^{\text{EMA}}$ takes the form*

$$\nabla_\theta \mathcal{L}_{\text{MF}}^{\text{EMA}} = 2 \mathbb{E}[(\nabla_\theta \mathbf{u}_\theta(\mathbf{x}_t, r, t))^\top \tilde{\mathbf{r}}^{\text{EMA}}],$$

966 where $\tilde{\mathbf{r}}^{\text{EMA}} := V_\theta^{\text{EMA}} - \mathbf{v}(\mathbf{x}_t, t)$ is the mean-field residual under the EMA tangent, and no $\mathbf{J} \mathbf{v}'$ noise
 967 term appears.

968 *Proof.* The EMA velocity anchor $\mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t)$ is deterministic given $\mathbf{x}_t$, since the EMA parameters $\bar{\theta}$ are fixed during the gradient step. Under stop-gradient, $V_{\bar{\theta}}^{\text{EMA}}$ depends on θ only through $\mathbf{u}_{\theta}(\mathbf{x}_t, r, t)$. By the same argument as equation 12, the per-step gradient is:

$$\nabla_{\theta} \ell = 2 (\nabla_{\theta} \mathbf{u}_{\theta})^{\top} (V_{\bar{\theta}}^{\text{EMA}} - \mathbf{v}_{\text{cond}}). \quad (19)$$

971 Writing $\mathbf{v}_{\text{cond}} = \mathbf{v}(\mathbf{x}_t, t) + \mathbf{v}'$ and noting that $\tilde{\mathbf{r}}^{\text{EMA}} := V_{\bar{\theta}}^{\text{EMA}} - \mathbf{v}(\mathbf{x}_t, t)$ is deterministic given $\mathbf{x}_t$:

$$\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t} [\nabla_{\theta} \ell] = 2 (\nabla_{\theta} \mathbf{u}_{\theta})^{\top} \tilde{\mathbf{r}}^{\text{EMA}}, \quad (20)$$

972 since $\mathbb{E}[\mathbf{v}' | \mathbf{x}_t] = \mathbf{0}$. The tower property $\nabla_{\theta} \mathcal{L}_{\text{MF}}^{\text{EMA}} = \mathbb{E}_{r, t, \mathbf{x}_t} [\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t} [\nabla_{\theta} \ell]]$ gives the stated result. Crucially, no $\mathbf{J}\mathbf{v}'$ term appears because the EMA tangent is deterministic; compare with the original MeanFlow gradient equation 12 where the stochastic tangent produces the $\mathbf{J}\mathbf{v}'$ noise. $\square$

975 E.4 Proof of Proposition 3 (Bias-Variance Tradeoff)

976 **Proposition 8** (Bias-Variance Tradeoff). *The EMA mean-field residual decomposes as*

$$\tilde{\mathbf{r}}^{\text{EMA}} = \mathbf{r}_{\theta} + (t-r) \partial_{\mathbf{z}} \mathbf{u}_{\theta}(\mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t) - \mathbf{v}(\mathbf{x}_t, t)),$$

977 where $\mathbf{r}_{\theta}$ is the true MeanFlow residual from equation 5. At convergence $\mathbf{r}_{\theta} \rightarrow \mathbf{0}$ and the FM anchor drives $\mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t) \rightarrow \mathbf{v}(\mathbf{x}_t, t)$, giving $\tilde{\mathbf{r}}^{\text{EMA}} \rightarrow \mathbf{0}$.

979 *Proof.* Expanding $V_{\bar{\theta}}^{\text{EMA}}$ from equation 9, the value of the compound prediction (ignoring the stop-gradient, which does not affect values) is:

$$V_{\bar{\theta}}^{\text{EMA}} = \mathbf{u}_{\theta} + (t-r)(\partial_{\mathbf{z}} \mathbf{u}_{\theta} \cdot \mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t) + \partial_t \mathbf{u}_{\theta}). \quad (21)$$

981 The true MeanFlow residual (with the marginal velocity as tangent) is:

$$\mathbf{r}_{\theta} = \mathbf{u}_{\theta} + (t-r)(\partial_{\mathbf{z}} \mathbf{u}_{\theta} \cdot \mathbf{v}(\mathbf{x}_t, t) + \partial_t \mathbf{u}_{\theta}) - \mathbf{v}(\mathbf{x}_t, t). \quad (22)$$

982 Subtracting:

$$\begin{aligned} \tilde{\mathbf{r}}^{\text{EMA}} &= V_{\bar{\theta}}^{\text{EMA}} - \mathbf{v}(\mathbf{x}_t, t) \\ &= \mathbf{r}_{\theta} + (t-r) \partial_{\mathbf{z}} \mathbf{u}_{\theta}(\mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t) - \mathbf{v}(\mathbf{x}_t, t)). \end{aligned} \quad (23)$$

983 At convergence, the true residual $\mathbf{r}_{\theta} \rightarrow \mathbf{0}$ (the MeanFlow identity is satisfied) and the FM anchor loss equation 10 supervises the boundary condition, driving $\mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t) \rightarrow \mathbf{v}(\mathbf{x}_t, t)$ and hence $\tilde{\mathbf{r}}^{\text{EMA}} \rightarrow \mathbf{0}$. $\square$

986 E.5 Proof of Proposition 4 (Target/Tangent Bias Asymmetry)

987 **Proposition 9** (Target/Tangent Bias Asymmetry). *Let $\hat{\mathbf{v}} = \mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t)$ be a deterministic proxy with bias $\mathbf{b} \triangleq \hat{\mathbf{v}} - \mathbf{v}(\mathbf{x}_t, t)$. For the tangent-replacement loss (the $\beta = 1$ corner), the loss minimum $\mathbf{u}^*$ satisfies $\mathbf{u}^*(\mathbf{x}_t, r, t) - \mathbf{v}(\mathbf{x}_t, t) = -\frac{1}{2}(t-r) \partial_{\mathbf{z}} \mathbf{u}^* \mathbf{b} + \mathcal{O}((t-r) \|\mathbf{b}\|^2)$. For the target-replacement loss, the loss minimum at $r=t$ satisfies $\mathbf{u}^*(\mathbf{x}_t, t, t) - \mathbf{v}(\mathbf{x}_t, t) = \mathbf{b}$.*

991 *Proof. Tangent replacement.* Substituting $\hat{\mathbf{v}}$ for $\mathbf{v}_{\text{cond}}$ only in the JVP tangent (target unchanged) gives the per-sample compound prediction

$$V_{\bar{\theta}}^{\text{EMA}} = \mathbf{u}_{\theta} + (t-r)(\partial_{\mathbf{z}} \mathbf{u}_{\theta} \cdot \hat{\mathbf{v}} + \partial_t \mathbf{u}_{\theta}),$$

993 and the loss is $\mathbb{E}[\|V_{\bar{\theta}}^{\text{EMA}} - \mathbf{v}_{\text{cond}}\|^2]$. By stationarity, $\mathbb{E}[V_{\bar{\theta}}^{\text{EMA}} - \mathbf{v}_{\text{cond}} | \mathbf{x}_t] = \mathbf{0}$ at the minimum. Taking conditional expectation and using $\mathbb{E}[\mathbf{v}_{\text{cond}} | \mathbf{x}_t] = \mathbf{v}(\mathbf{x}_t, t)$,

$$\mathbf{u}^* + (t-r)(\partial_{\mathbf{z}} \mathbf{u}^* \cdot \hat{\mathbf{v}} + \partial_t \mathbf{u}^*) = \mathbf{v}(\mathbf{x}_t, t).$$

995 At the same point, the *true* MeanFlow identity (with the marginal velocity as tangent) reads

$$\mathbf{u}^{\diamond} + (t-r)(\partial_{\mathbf{z}} \mathbf{u}^{\diamond} \cdot \mathbf{v} + \partial_t \mathbf{u}^{\diamond}) = \mathbf{v}(\mathbf{x}_t, t),$$

996 with $\mathbf{u}^{\diamond}(\mathbf{x}_t, t, t) = \mathbf{v}(\mathbf{x}_t, t)$. Subtracting and writing $\mathbf{u}^* = \mathbf{u}^{\diamond} + \Delta$,

$$\Delta + (t-r) \partial_{\mathbf{z}} \mathbf{u}^* \mathbf{b} + (t-r) \text{D}_t \Delta = \mathbf{0}, \quad \text{D}_t \Delta \triangleq \partial_t \Delta + \mathbf{v} \cdot \partial_{\mathbf{z}} \Delta,$$

997 with the boundary condition $\Delta(\mathbf{x}_t, t, t) = \mathbf{0}$. Substituting $s = t - r$ at fixed r converts this to a
 998 first-order linear ODE in s :

$$s \frac{d\Delta}{ds} + \Delta = -s \partial_{\mathbf{z}} \mathbf{u}^* \mathbf{b} + \mathcal{O}(\|\mathbf{b}\|^2),$$

999 which has $s\Delta(s) = -\frac{s^2}{2} \partial_{\mathbf{z}} \mathbf{u}^* \mathbf{b} + C$; the boundary condition $\Delta(0) = \mathbf{0}$ forces $C = 0$. Hence

$$\Delta(\mathbf{x}_t, r, t) = -\frac{1}{2}(t - r) \partial_{\mathbf{z}} \mathbf{u}^* \mathbf{b} + \mathcal{O}((t - r) \|\mathbf{b}\|^2),$$

1000 i.e., the stationary error is $\mathcal{O}((t - r) \|\partial_{\mathbf{z}} \mathbf{u}\| \|\mathbf{b}\|)$ and vanishes at the boundary $r \rightarrow t$. The factor of $\frac{1}{2}$
 1001 comes from integrating the material-derivative term, which contributes equally to the leading-order
 1002 coefficient as the bias source term.

1003 **Target replacement.** Substituting $\hat{\mathbf{v}}$ for $\mathbf{v}_{\text{cond}}$ in the regression target yields the loss $\mathbb{E}\|\mathbf{u}_{\theta} + (t -$
 1004 $r) D\mathbf{u}/Dt - \hat{\mathbf{v}}\|^2$. By stationarity, the loss minimum satisfies $\mathbb{E}[\mathbf{u}^* + (t - r) D\mathbf{u}^*/Dt - \hat{\mathbf{v}} \mid \mathbf{x}_t] = \mathbf{0}$.
 1005 Specializing to $r = t$, the gap term vanishes and the equation collapses to $\mathbf{u}^*(\mathbf{x}_t, t, t) = \hat{\mathbf{v}} = \mathbf{v}(\mathbf{x}_t, t) +$
 1006 $\mathbf{b}$, so the stationary error at the boundary is exactly $\mathbf{b}$, independent of (t, r) .

1007 **Asymmetry.** Tangent replacement carries a multiplicative $(t - r)$ factor on the bias and is therefore
 1008 pinned to zero at the boundary $r = t$ that the FM anchor enforces; target replacement carries no such
 1009 factor and the bias persists at every t , requiring a boundary regularizer at *every* (r, t) to remove. $\square$

1010 E.6 Proof of Theorem 4 (Optimal Control-Variate Coefficient)

1011 **Theorem 10** (Optimal control-variate coefficient). *Under the scalar-isotropic approximation $\Sigma_{\mathbf{v}'} \approx$*
 1012 *$\sigma^2 \mathbf{I}_d$ and $\mathbf{J} \approx \kappa \mathbf{I}_d$, the conditional MSE of $\mathbf{g}^{(\beta)}$ in equation 7 is*

$$M(\beta) = \beta^2 (\kappa + 1)^2 \|\mathbf{b}\|^2 + \sigma^2 d ((1 - \beta)\kappa - \beta)^2,$$

1013 with unique minimizer $\beta^* = \frac{\kappa}{\kappa + 1} \cdot \frac{\sigma^2 d}{\sigma^2 d + \|\mathbf{b}\|^2}$ and optimum value $M(\beta^*) = \sigma^2 d \kappa^2 \|\mathbf{b}\|^2 / (\sigma^2 d + \|\mathbf{b}\|^2)$.
 1014 For all $\kappa, \|\mathbf{b}\|^2 > 0$, the optimum strictly dominates both $M(0) = \sigma^2 \kappa^2 d$ and $M(1) = (\kappa + 1)^2 \|\mathbf{b}\|^2 +$
 1015 $\sigma^2 d$.

1016 *Proof. Step 1: bias-variance decomposition.* From equation 7,

$$\mathbf{g}^{(\beta)} = 2\mathbf{g}^\top \left[\bar{\mathbf{r}} + ((1 - \beta)\mathbf{J} - \beta\mathbf{I}) \mathbf{v}' + \beta(\mathbf{J} + \mathbf{I}) \mathbf{b} \right].$$

1017 Using $\mathbb{E}_{\mathbf{x}_0 | \mathbf{x}_t}[\mathbf{v}'] = \mathbf{0}$, the conditional mean is $\mathbb{E}[\mathbf{g}^{(\beta)} \mid \mathbf{x}_t] = 2\mathbf{g}^\top (\bar{\mathbf{r}} + \beta(\mathbf{J} + \mathbf{I}) \mathbf{b})$ and the centered
 1018 noise is $-2\mathbf{g}^\top ((1 - \beta)\mathbf{J} - \beta\mathbf{I}) \mathbf{v}'$ (relative to the ideal target $2\mathbf{g}^\top \bar{\mathbf{r}}$). The conditional MSE w.r.t. the
 1019 ideal target therefore decomposes as

$$\text{MSE} = \underbrace{4\beta^2 \|\mathbf{g}^\top (\mathbf{J} + \mathbf{I}) \mathbf{b}\|^2}_{\text{bias}^2} + \underbrace{4 \text{Tr}(\mathbf{g} \mathbf{g}^\top ((1 - \beta)\mathbf{J} - \beta\mathbf{I}) \Sigma_{\mathbf{v}'} ((1 - \beta)\mathbf{J} - \beta\mathbf{I})^\top)}_{\text{variance}}.$$

1020 **Step 2: matrix-form expansion and strict convexity.** Let $\mathbf{A} \triangleq \mathbf{J} + \mathbf{I}$ and $\mathbf{G}_\theta \triangleq \mathbf{g} \mathbf{g}^\top \succeq \mathbf{0}$. Substi-
 1021 tuting $((1 - \beta)\mathbf{J} - \beta\mathbf{I}) = \mathbf{J} - \beta\mathbf{A}$ and using $\|\mathbf{g}^\top \mathbf{A} \mathbf{b}\|^2 = \text{Tr}(\mathbf{G}_\theta \mathbf{A} \mathbf{b} \mathbf{b}^\top \mathbf{A}^\top)$, the conditional MSE
 1022 writes as a quadratic in β :

$$M(\beta) = \alpha_0 - 2\alpha_1 \beta + \alpha_2 \beta^2, \tag{24}$$

1023 with (suppressing the constant prefactor 4)

$$\begin{aligned} \alpha_0 &\triangleq \text{Tr}(\mathbf{G}_\theta \mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{J}^\top), \\ \alpha_1 &\triangleq \text{Tr}(\mathbf{G}_\theta \mathbf{J} \Sigma_{\mathbf{v}'} \mathbf{A}^\top), \\ \alpha_2 &\triangleq \text{Tr}(\mathbf{G}_\theta \mathbf{A} (\Sigma_{\mathbf{v}'} + \mathbf{b} \mathbf{b}^\top) \mathbf{A}^\top). \end{aligned}$$

1024 (We used the symmetry of $\mathbf{G}_\theta$ and $\Sigma_{\mathbf{v}'}$ to consolidate the linear-in- β trace into a single α_1 .) The
 1025 Hessian is

$$\frac{\partial^2 M}{\partial \beta^2} = 2\alpha_2 = 2 \text{Tr}(\mathbf{G}_\theta \mathbf{A} (\Sigma_{\mathbf{v}'} + \mathbf{b} \mathbf{b}^\top) \mathbf{A}^\top) \geq 0,$$

1026 since each factor inside the trace is positive semi-definite ($\mathbf{G}_\theta \succeq \mathbf{0}$, $\Sigma_{\mathbf{v}'} + \mathbf{b} \mathbf{b}^\top \succeq \mathbf{0}$, and $\mathbf{A}(\Sigma_{\mathbf{v}'} +$
 1027 $\mathbf{b} \mathbf{b}^\top) \mathbf{A}^\top = \mathbf{M} \mathbf{M}^\top \succeq \mathbf{0}$ for $\mathbf{M} \triangleq \mathbf{A}(\Sigma_{\mathbf{v}'} + \mathbf{b} \mathbf{b}^\top)^{1/2}$). Hence $M(\beta)$ is convex in β . The Hessian is

1028 strictly positive whenever $\mathbf{G}_\theta^{1/2} \mathbf{M} \neq \mathbf{0}$, which holds outside two pathological regimes: (i) $\mathbf{A} = \mathbf{0}$,
 1029 equivalently $\mathbf{J} = -\mathbf{I}$ (network with $\partial_z \mathbf{u}_\theta \equiv \mathbf{0}$, e.g., at saturated initialization), in which case the
 1030 linear coefficient α_1 also vanishes and $M(\beta)$ is constant; (ii) $\mathbf{g} = \mathbf{0}$ (frozen parameters). Outside
 1031 these, $M(\beta)$ is *strictly* convex with unique unconstrained minimizer

$$\beta_{\text{matrix}}^* = \frac{\alpha_1}{\alpha_2} = \frac{\text{Tr}(\mathbf{G}_\theta \mathbf{J} \Sigma_{\mathbf{v}'} (\mathbf{J} + \mathbf{I})^\top)}{\text{Tr}(\mathbf{G}_\theta (\mathbf{J} + \mathbf{I}) (\Sigma_{\mathbf{v}'} + \mathbf{b} \mathbf{b}^\top) (\mathbf{J} + \mathbf{I})^\top)}. \quad (25)$$

1032 Note that β_{matrix}^* may fall outside $[0, 1]$; the box-constrained optimum on the $\beta \in [0, 1]$ interval
 1033 (cf. equation 6) is then $\text{clip}(\beta_{\text{matrix}}^*, 0, 1)$, attained at the nearer endpoint by convexity.

1034 **Step 3: scalar approximation.** Substituting $\mathbf{J} = \kappa \mathbf{I}_d$, $\Sigma_{\mathbf{v}'} = \sigma^2 \mathbf{I}_d$, $\mathbf{A} = (\kappa + 1) \mathbf{I}_d$, and the
 1035 parameter-isotropy approximation $\mathbf{G}_\theta \propto \mathbf{I}_d$ (which absorbs the parameter-direction normalization
 1036 into a per-component MSE),

$$\alpha_0 \propto \sigma^2 \kappa^2 d, \quad \alpha_1 \propto \sigma^2 \kappa (\kappa + 1) d, \quad \alpha_2 \propto (\kappa + 1)^2 (\sigma^2 d + \|\mathbf{b}\|^2),$$

1037 which equivalently gives

$$M(\beta) \propto \beta^2 (\kappa + 1)^2 \|\mathbf{b}\|^2 + \sigma^2 d ((1 - \beta) \kappa - \beta)^2.$$

1038 **Step 4: minimization.** Differentiating and setting to zero (or directly evaluating α_1/α_2),

$$\beta^* = \frac{\alpha_1}{\alpha_2} = \frac{\sigma^2 \kappa (\kappa + 1) d}{(\kappa + 1)^2 (\sigma^2 d + \|\mathbf{b}\|^2)} = \frac{\kappa}{\kappa + 1} \cdot \frac{\sigma^2 d}{\sigma^2 d + \|\mathbf{b}\|^2}.$$

1039 **Step 5: optimum value and corner dominance.** At β^* , $(1 - \beta^*) \kappa - \beta^* = \kappa \|\mathbf{b}\|^2 / (\sigma^2 d + \|\mathbf{b}\|^2)$.
 1040 Substituting back,

$$M(\beta^*) \propto \frac{\sigma^2 d \kappa^2 \|\mathbf{b}\|^2}{\sigma^2 d + \|\mathbf{b}\|^2}.$$

1041 For $\kappa > 0$ and $\|\mathbf{b}\|^2 > 0$: $M(\beta^*)/M(0) = \|\mathbf{b}\|^2 / (\sigma^2 d + \|\mathbf{b}\|^2) < 1$, so $M(\beta^*) < M(0)$. For the $M(1)$
 1042 comparison, $M(1) - M(\beta^*) \propto (\kappa + 1)^2 \|\mathbf{b}\|^2 + \sigma^2 d - \sigma^2 d \kappa^2 \|\mathbf{b}\|^2 / (\sigma^2 d + \|\mathbf{b}\|^2)$, and the first two
 1043 terms exceed $\sigma^2 d \kappa^2 \|\mathbf{b}\|^2 / (\sigma^2 d + \|\mathbf{b}\|^2)$ since $(\kappa + 1)^2 > \kappa^2$ for $\kappa > 0$. Strict dominance is therefore
 1044 over the unconstrained interior optimum; if $\beta_{\text{matrix}}^* \notin [0, 1]$ the box-constrained optimum is at the
 1045 nearer corner and the dominance is non-strict. $\square$

1046 F DiT Latent-Space Training on ImageNet

1047 We test whether the variance-reduction mechanism of Theorems 2 and 4 scales beyond toys by
 1048 training DiT [42] variants on ImageNet in the latent space of a frozen Stable Diffusion VAE, follow-
 1049 ing the recipe of [21]. The four β configurations share identical hyperparameters and code paths;
 1050 only the loss differs (vanilla MeanFlow MSE for the interior β runs versus the EMA-tangent loss
 1051 in equation 11 for the $\beta = 1$ corner).

1052 F.1 DiT-B/4, ImageNet-256 $\times$ 256

1053 Setup: logit-normal (r, t) sampling with mean -0.4 and stddev 1, overlap rate 0.75, AdamW with
 1054 constant learning rate 10^{-4} , EMA decay 0.9999, batch size 256 across 32 TPU v4 chips. All four
 1055 configurations ($\beta \in \{0, 0.25, 0.5, 1\}$) share the same DiT-B/4 backbone, data pipeline, optimizer,
 1056 EMA schedule, and (r, t) sampling distribution. The interior β runs ($\beta \in \{0, 0.25, 0.5\}$) reuse the
 1057 baseline configuration verbatim and add only the tangent-mixing rule of equation 6; the $\beta = 1$ corner
 1058 additionally disables the Karras-style adaptive loss weighting and adds the FM anchor of equation 10
 1059 ($\delta \sim \mathcal{U}[0, 10^{-3}]$, $\lambda = 0.1$). The extra EMA forward pass and FM-anchor evaluation per step add up to
 1060 a $\sim 22\%$ wall-clock overhead on TPU v4 (2.20 steps/s for $\beta = 1$ versus 2.65 steps/s for the baseline);
 1061 the $\beta \in \{0.25, 0.5\}$ interior runs incur only the EMA forward overhead ($\sim 10\%$).

1062 **Three-floor structure and matched-step ordering.** Table 5 reports FID_{50k} at representative 5k-
 1063 step checkpoints for all three runs, and Figure 5 renders the full trajectories on a semilog-y axis.
 1064 After the early-training transient ($\leq 10k$ steps), the empirical ordering $\text{FID}(\beta=0) < \text{FID}(\beta=0.5) <$

1065 FID($\beta=1$) holds at *every* matched-step checkpoint we logged (24 consecutive datapoints from step
1066 30k onward). The trajectory has two phases: in the first (steps 20k to 45k), the gap between $\beta=0.5$
1067 and the $\beta=1$ corner widens monotonically (from +2.1 at step 20k to -11.4 at step 45k), validating
1068 the interior optimum predicted by Theorem 4; once the gradient-noise reduction of $\beta=0.5$ has paid
1069 off, the convergence-rate advantage stabilizes. In the second phase (≥ 90 k), the gap between $\beta=0.5$
1070 and the $\beta=0$ baseline reaches an asymptotic offset of $\boxed{+7.77 \pm 0.11}$ FID over 11 consecutive
1071 matched-step windows (steps 95k to 145k), well below the +16 FID gap of the early-training phase.

1072 **$\beta=0.5$ beats $\beta=1$'s asymptotic floor.** The $\beta=1$ corner run completed its scheduled 300k-step
1073 training with a final logged FID of **23.36** (step 295k); we treat this as the empirical $\beta=1$ floor.
1074 The $\beta=0.5$ run also completed its scheduled 300k steps with a final logged FID of **12.51** (step
1075 295k), a margin of 10.85 FID below the $\beta=1$ asymptote. The four converged FID floors recover
1076 the bias-variance ordering of Theorem 4:

$$\text{FID}^*(\beta=0) \approx 10.34 < \text{FID}^*(\beta=0.25) \lesssim 14.79 < \text{FID}^*(\beta=0.5) = 12.51 < \text{FID}^*(\beta=1) = 23.36.$$

1077 (The baseline minimum FID of 10.34 was logged at step 145k; subsequent training showed mild
1078 late-stage drift to 12.12 at step 235k, so the asymptotic baseline floor lies in $[10.3, 12.3]$.) The
1079 $\beta=0.25$ run is still in training; its latest logged FID at step 170k is 14.79 and remains a strict upper
1080 bound on its own asymptote.

1081 Quantitatively, taking offsets from the best-logged baseline floor 10.34, $\Delta\text{FID}(\beta=0.5) \approx 2.17$
1082 and $\Delta\text{FID}(\beta=1) \approx 13.02$, giving an empirical FID-offset ratio of $\sim 6:1$. The gradient-MSE bias
1083 term $\beta^2(\kappa+1)^2\|\mathbf{b}\|^2$ predicts the much smaller 1:4 ratio (from β^2 alone). The FID landscape is
1084 therefore *super-linear* in the MSE-axis bias: it penalizes large bias more aggressively than gradient
1085 MSE does, which pulls the FID-optimal β past the gradient-MSE interior minimizer to the unbiased
1086 corner $\beta=0$. The four-floor structure is the strongest empirical confirmation of the bias-variance
1087 decomposition we obtain at DiT scale.

Table 5: DiT-B/4 / ImageNet-256 FID_{50k} (lower is better) at representative 5k-step checkpoints across the four-point β -sweep. The baseline ($\beta=0$) and $\beta=1$ corner runs were trained to completion (235k and 295k logged steps, respectively); the interior probe at $\beta=0.5$ completed its scheduled 300k steps with final FID 12.51 at step 295k; the interior probe at $\beta=0.25$ is still training at the time of submission, with the latest logged FID at step 145k. Bottom rows give the matched-step deltas. The four-point ordering $\text{FID}(\beta=0) < \text{FID}(\beta=0.25) < \text{FID}(\beta=0.5) < \text{FID}(\beta=1)$ holds at every matched-step checkpoint.

step (k)	30	40	50	60	80	100	115	130	145	final
baseline ($\beta=0$)	128.9	79.2	54.5	41.8	28.3	17.7	14.4	12.1	10.3	12.1 [†]
$\beta=0.25$	138.5	88.4	61.6	46.7	30.4	23.3	20.3	18.3	16.7	(running) [‡]
$\beta=0.5$	145.2	94.6	n/a	n/a	n/a	25.5	22.2	19.7	18.1	12.5
$\beta=1$	148.8	104.1	79.9	65.7	49.9	41.7	38.0	35.2	32.7	23.4
$\Delta(\beta=0.25 - \text{baseline})$	+9.6	+9.2	+7.1	+4.9	+2.1	+5.6	+5.9	+6.2	+6.4	n/a
$\Delta(\beta=0.5 - \text{baseline})$	+16.3	+15.4	n/a	n/a	n/a	+7.8	+7.8	+7.6	+7.8	n/a
$\Delta(\beta=1 - \text{baseline})$	+19.9	+24.9	+25.4	+23.9	+21.6	+24.0	+23.6	+23.1	+22.4	n/a

[†] The baseline run was stopped at step 235k with last logged FID 12.12; its minimum logged FID was 10.34 at step 145k, with mild late-training drift thereafter. We report the late-training value as the empirical asymptotic floor for the matched-step comparison.

[‡] Latest logged $\beta=0.25$ FID at the time of submission is 14.79 at step 170k, still descending; the run is scheduled for 300k total steps.

1088 **Bias and shrinkage trajectory.** To localize the predicted optimum we measure two diagnos-
1089 tics on the DiT checkpoints: an EMA-tracking proxy $\|\widehat{\mathbf{b}}\|^2 \triangleq \mathbb{E}_{\mathbf{x}_t} [\|\mathbf{u}_\theta(\mathbf{x}_t, t, t) - \mathbf{u}_{\bar{\theta}}(\mathbf{x}_t, t, t)\|^2]$
1090 (the boundary gap between current parameters and the EMA copy), and the irreducible noise floor
1091 $\sigma^2 d \approx 8.2 \times 10^3$ from $\text{Tr}(\Sigma_{v'})$ on the same batch. Table 6 reports the trajectory: the EMA-tracking
1092 proxy *decays* through training, and the shrinkage factor $\sigma^2 d / (\sigma^2 d + \|\widehat{\mathbf{b}}\|^2)$ stays near 1 through-
1093 out converged training. We caveat that $\|\widehat{\mathbf{b}}\|^2$ is the EMA-tracking gap, not the true model bias
1094 $\|\mathbf{u}_\theta(\mathbf{x}_t, t, t) - \mathbf{v}(\mathbf{x}_t, t)\|^2$; cleanly estimating the latter requires access to the conditional distribu-
1095 tion $p(\mathbf{x}_0 | \mathbf{x}_t)$, which is intractable on ImageNet. Combined with $\kappa \approx 1$ from the Frobenius-norm
1096 Hutchinson trace, the scalar-isotropic substitution gives $\beta^* \in [0.46, 0.50]$ across the full trajectory;

1097 this is a coarse first-cut prediction that the direct matrix-form measurement (Section 5.4) refines to
1098 $\beta_{\text{no bias}}^* \approx 0.94$.

Table 6: Direct DiT-checkpoint measurement of the bias-variance decomposition ingredients of Theorem 4 at $t = 0.5$. The EMA bias $\|\mathbf{b}\|^2$ averages over 128 synthetic-Gaussian $\mathbf{x}_t$ samples and is reported for both methods (which track within 20% of each other). The shrinkage factor uses $\sigma^2 d \approx 8.2 \times 10^3$. The predicted β^* uses $\kappa = 1$ from the Hutchinson trace.

step (k)	$\ \mathbf{b}\ ^2$ (avg)	$\sigma^2 d / (\sigma^2 d + \ \mathbf{b}\ ^2)$	β^*
20	~ 700	0.92	0.46
40	~ 150	0.98	0.49
80	~ 50	0.99	0.50